

Cassandra User Group Cologne

19:00 - Reception


19:20 - Cassandra libraries for Java developers

DuyHai Doan, Cassandra Evangelist at DataStax

20:20 - Apache Cassandra 3.0

Robert Stupp, Committer to Apache Cassandra, CIO contentteam AG

21:00 - Finish & Networking





contentteam

APACHE CASSANDRA 3.0

CASSANDRA USER GROUP COLOGNE

24.03.2015

Robert Stupp



- CIO contentteam
- Committer to Apache Cassandra
- Coding experience since 1985
- Internet and related technologies since 1992

- rstupp@contentteam.com
- @snazy



contentteam is a DataStax Solutions Partner

contentteam is active in Apache Cassandra community

1. Cassandra history (short)
2. Cassandra 3.0
3. Cassandra community
4. Apache Cassandra vs.
DataStax Enterprise
5. One more thing :)

APACHE CASSANDRA HISTORY

- Initially developed at Facebook to build a “continuously available” database
 - Replication
 - Globally distributed
 - Masterless architecture
- Influenced by BigTable and Dynamo
- Today:
 - Huge amount of working installation - few nodes up to 1000+ globally distributed nodes

- Open sourced in 2008
- Version 0.3 - July 2009
- Version 0.6 - June 2010
- Version 0.7 - 2011
- Version 1.0 - October 2011 - introduction of DSE
- Version 1.2 - December 2012
- Version 2.0 - August 2013
- Version 2.1 - September 2014

- Version 2.0 line
critical bugfixes applied to 2.0 release
- Version 2.1 line
bugfixes applied to 2.0 release
some new, non-intrusive features
- Version 3.0
lots of new features
lots of improvements

APACHE CASSANDRA 3.0

DISCLAIMER

- Apache Cassandra 3.0 is **still in development**
- Features might be changed / revoked until 3.0 release

RELEASE DATE

- Apache Cassandra 3.0 will be released **when it is finished**
- Don't ask for a release date - we don't know it yet ;)

CASSANDRA 3.0 FEATURES

- JSON support
- User-Defined-Functions + User-Defined-Aggregates
- Role based access control
- New row cache
- Lots of (small) performance improvements summing up to a huge improvement
- Altogether approx >100 tickets for 3.0
- Plus changes merged from 2.0 via 2.1 to 3.0 and 2.1 to 3.0

- `cassandra-cli` is removed (as announced)

JSON Support

CASSANDRA 3.0

- Allows you to do **INSERT** and **SELECT** data using JSON data format
- Format your data to insert using JSON - can save a step to transform data
- **Ease web application development** - less transformation to/from Cassandra to the browser
- Also nice when using NodeJS

```
CREATE TYPE address (  
  street text,  
  city text,  
  zip_code int,  
  phones set<text>  
);
```

```
CREATE TABLE users (  
  id uuid PRIMARY KEY,  
  name text,  
  addresses frozen<map<text, address>>  
);
```

this is JSON

```
INSERT INTO users JSON  
'{"id": "4b856557-7153",  
  "name": "snazy",  
  "addresses": {"work": {"street": "Im Mediapark 6",  
    "city": "Köln",  
    "zip_code": 50670,  
    "phones": ["+492214546200"]}}}'
```

```
cqlsh> SELECT JSON * FROM users;
```

```
[json]
```

```
{ "id": "4b856557-7153",  
  "name": "snazy",  
  "addresses": { "work": { "street": "Im Mediapark 6",  
                           "city": "Köln",  
                           "zip_code": 50670,  
                           "phones": [ "+492214546200" ] } } }
```

```
( 1 rows )
```

this is JSON

- JSON support **does not** introduce schema-free tables!
- <http://rustyrazorblade.com/2014/07/the-myth-of-schema-less/>
- <https://blog.compose.io/schema-less-is-usually-a-lie/>

User-Defined-Functions (UDFs)

User-Defined-Aggregates (UDAs)

CASSANDRA 3.0

- UDF means User Defined Function
- **You** write the code that's executed on Cassandra nodes
- Functions are **distributed** transparently to the whole cluster
- *You may not have to wait for a new release for new functionality :)*

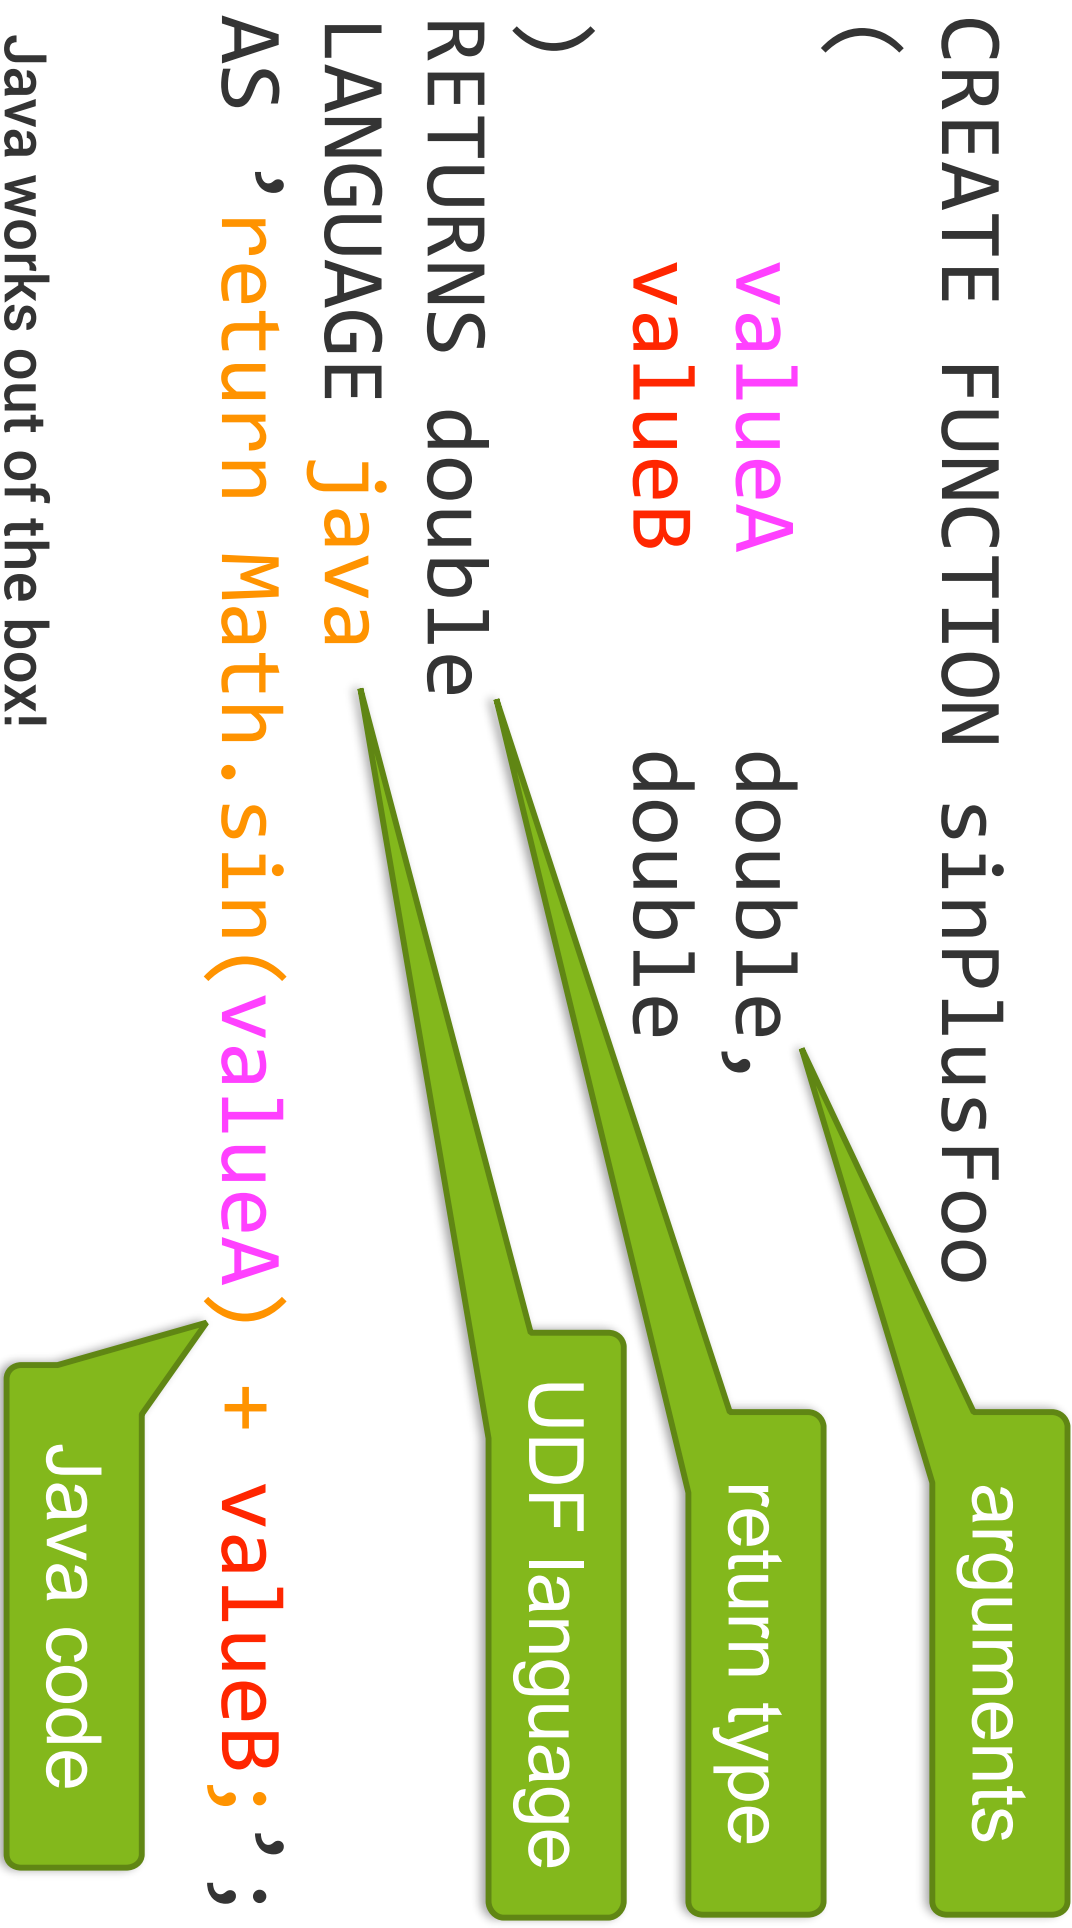
- „Pure“
 - just input parameters
 - no state, side effects, dependencies to other code, etc
- Usually deterministic

Consider a Java function like...

```
import nothing;  
  
public final class MyClass  
{  
    public static int myFunction ( int argument )  
    {  
        return argument * 42;  
    }  
}
```

This would be your UDF

```
CREATE FUNCTION sinPlusFoo  
(  
    valueA double,  
    valueB double  
)  
RETURNS double  
LANGUAGE java  
AS 'return Math.sin(valueA) + valueB;';
```



Java works out of the box!

```
CREATE FUNCTION sin (
    value double )
RETURNS double
LANGUAGE javascript
AS 'Math.sin(value);';
```

JavaScript
works, too

JavaScript code

JavaScript works out of the box!

- “Scripting for the Java Platform”
- UDFs can be written in **Java** and **JavaScript**
- Optionally: **Groovy**, **JRuby**, **Jython**, **Scala**
- Not: Clojure (JSR 223 implementation’s wrong)

- Builds Java (or script) source
- Compiles that code (Java class, or compiled script)
- Loads the compiled code
- Migrates the function to all other nodes
- Done - UDF is executable on any node

- Support for **all** Cassandra **types** for arguments and return value
- All means
 - Primitives (**boolean, int, double, uuid, etc**)
 - Collections (**list, set, map**)
 - Tuple types, User Defined Types

UDF - For what?

```
SELECT sumThat ( colA, colB )  
FROM myTable  
WHERE key = ...
```

```
SELECT sin ( foo )  
FROM myCircle  
WHERE pk = ...
```

Now your application can sum two values
in one row - or create the sin of a value!

GREAT NEW FEATURES!

Okay - not really...

- UDFs on their own are just „nice to have“
- Nothing you couldn't do better in your application

User Defined Aggregates !

Use UDFs to code your own aggregation functions

(Aggregates are things like SUM, AVG, MIN, MAX, etc)

Aggregates :

consume values from **multiple rows** &
produce a **single result**

```
CREATE AGGREGATE minimum ( int )
```

arguments

```
STYPE int  
SFUNC minimumState;
```

state type

name of the
state UDF

```
SELECT minimum ( val ) FROM foo
```

1. **Initial state** is set to **null**
2. for **each row** the **state function** is called with current state and column value - returns **new state**
3. After all rows the **aggregate returns** the **last state**

```
CREATE AGGREGATE average ( int )
SFUNC averageState
STYPE tuple<long,int>
FINALFUNC averageFinal
INITCOND ( 0, 0 );
```

initial state
value

name of the
final UDF

```
SELECT average ( val ) FROM foo ...
```

1. **Initial state** is set to $(0,0)$
2. for **each row** the **state function** is called with
current state + column value - returns **new state**
3. After all rows the **final function** is called with **last state**
4. **final function** calculates the **aggregate**

**Now everybody
can execute
evil
code on your
cluster :)**

- There will be permissions to restrict (allow)
- UDF creation (DDL)
- UDF execution (DML)

Keep in mind:

- JSR-223 has overhead - Java UDFs are much faster
- Do not allow everyone to create UDFs (in production)
- Keep your UDFs “**pure**”
- **Test** your UDFs and user defined aggregates **thoroughly**

- **UDFs** and user defined aggregates are **executed** on the **coordinator node**
- **Prefer** to use **Java-UDFs** for performance reasons

UDFs could be useful for...

- Functional indexes
- Partial indexes
- Filtering
- Distributed GROUP BY
- etc etc

NOT IN C*
3.0 !

Role based access control

CASSANDRA 3.0

- Grant/revoke permissions to/from roles
- Grant roles to users

```
GRANT <permission> ON <resource>  
TO [[USER] <username> | ROLE <rolename>]
```

```
REVOKE <permission> ON <resource>  
FROM [[USER] <username> | ROLE <rolename>]
```

```
LIST <permissionOrAll> [ON <resource>]  
[OF [[USER] <username> | ROLE <rolename>]  
[NORECURSIVE]
```

- Authentication/authorization has been reworked for Cassandra 3.0
- Much more options and possibilities
- See CASSANDRA-8394 for more information

New row cache

CASSANDRA 3.0

- Old row cache recommendation: "don't use it"
- Old row cache had data in off-heap, but management data in Java heap
- Resulted in a lot of additional GC pressure

- All data (whole concurrent hash map) is off-heap
- Uses APL2 licensed <https://github.com/snazy/ohc>
- Works with really big row cache
- Works on really big machines
- But don't expect a huge performance improvement
- Serialization of data to/from off-heap is still a bottleneck
- Will work on that bottleneck in future versions

More improvements

CASSANDRA 3.0

- Refactor and modernize storage engine (CASSANDRA-8099)
- Modernize schema tables (CASSANDRA-6717)
- CQL row read optimization
- Reduce GC pressure
- Make internal nomenclature intuitive

Memory related

- Support direct buffer decompression for reads (CASSANDRA-8464)
- Avoid memory allocation when searching index summary (CASSANDRA-8793)
- Use preloaded jemalloc w/ Unsafe (CASSANDRA-8714)

Throughput/CPU related

- Improve concurrency of repair (CASSANDRA-6455, 8208)
- Select optimal CRC32 implementation at runtime (CASSANDRA-8614)
- plus many more

Windows

- Cassandra 3.0 is tested on Windows
- 3.0 works definitely better on Windows than 2.1
- But still some issues to solve

My personal recommendation

- Use Cassandra on Linux

CASSANDRA >= 3.1 FEATURES

- Global indexes
- More UDF related stuff
 - Function based indexes
 - Use UDFs in filtering clauses
 - Distributed aggregates
- RAMP transactions
- More internal improvements and optimizations
- Expect Thrift to disappear
- Remember:
Everything is subject to change if not released ;)

CASSANDRA COMMUNITY

- People writing code (of course ;))
- People doing talks and presentations
- People active on social media (Twitter, LinkedIn, SlideShare, YouTube)
- People active on mailing list
- People working in the background
- **AND YOU !**

- Use the material provided by
 - DataStax
 - Planet Cassandra
 - “the community”
- on/via
 - YouTube
 - SlideShare
 - DataStax academy
 - Webinars
 - Meetups

APACHE CASSANDRA "VS." DATASTAX ENTERPRISE

- **Apache Cassandra** is open-source
- Support via user mailing list and tickets
- No commercial support
- **DataStax Enterprise** uses Apache Cassandra
- Adds graph database
- Adds enhanced security
- Adds analytics (Spark + Hadoop)
- Adds search (SolR)
- Adds commercial support

ONE
MORE
THING...



CASSANDRA-ON-MESOS

Program against your **datacenter** like it's a **single pool of resources**

Apache Mesos abstracts CPU, memory, storage, and other compute resources away from machines (physical or virtual), enabling fault-tolerant and elastic distributed systems to easily be built and run effectively.



— Your Apps

Deploy any Linux application on the Mesosphere DCOS with no code changes. Run your traditional applications, new applications and big data workloads with ease.



— Datacenter Services



Services such as Hadoop, Spark, Kafka, YARN and Kubernetes are the killer apps of the DCOS and can be installed with a single command.



— Mesosphere DCOS

The Mesosphere DCOS organizes the machines in your cluster. It provides an API for building and orchestrating distributed systems and a user interface to easily manage thousands of nodes.

- Run killer applications and services like Spark, Kafka and Cassandra
- In your own data center
- On Amazon EC2
- On Google GCE

- Allows to run Apache Cassandra on Apache Mesos
- Developed by  mesosphere and  contentteam
- Allows to spawn your Cassandra Cluster with a single command on Mesos
- Expect a first release-candidate this week !

Q & A

Robert Stupp

rstupp@contentteam.com
[@snazy](https://twitter.com/snazy)