



# User-Defined Functions & Materialized Views

DuyHai DOAN



# User Define Functions

- Why ?
- Detailed Impl
- UDAs
- Gotchas

# Rationale

- Push computation server-side
  - save network bandwidth (1000 nodes!)
  - simplify client-side code
  - provide standard & useful function (sum, avg ...)
  - accelerate analytics use-case (**pre-aggregation** for Spark)

# How to create an UDF ?

```
CREATE [OR REPLACE] FUNCTION [IF NOT EXISTS]
[keyspace.]functionName (param1 type1, param2 type2, ...)
CALL ON NULL INPUT | RETURNS NULL ON NULL INPUT
RETURN returnType
LANGUAGE language
AS '
    // source code here
';
```

# How to create an UDF ?

CREATE [OR REPLACE] FUNCTION [IF NOT EXISTS]

[keyspace.]functionName (param<sub>1</sub> type<sub>1</sub>, param<sub>2</sub> type<sub>2</sub>, ...)

CALLED ON NULL INPUT | RETURNS NULL ON NULL INPUT

RETURN returnType

LANGUAGE language

AS '

// source code here

';

CREATE OR REPLACE FUNCTION IF NOT EXISTS

# How to create an UDF ?

CREATE [OR REPLACE] FUNCTION [IF NOT EXISTS]

[*keyspace.*]***functionName*** (*param<sub>1</sub>* *type<sub>1</sub>*, *param<sub>2</sub>* *type<sub>2</sub>*, ...)

CALLED ON NULL INPUT | RETURNS NULL ON NULL INPUT

RETURN *returnType*

LANGUAGE

An UDF is *keyspace-wide*

AS '

// source code here

';

# How to create an UDF ?

```
CREATE [OR REPLACE] FUNCTION [IF NOT EXISTS]
[keyspace.]functionName (param1, type1, param2, type2, ...)
CALLED ON NULL INPUT | RETURNS NULL ON NULL INPUT
RETURN returnType
LANGUAGE language
AS '
    // source code here
';
```

Param name to refer to in the code  
Type = CQL3 type

# How to create an UDF ?

```
CREATE [OR REPLACE] FUNCTION [IF NOT EXISTS]
[keyspace.]functionName (param1 type1, param2 type2, ...)
CALLED ON NULL INPUT | RETURNS NULL ON NULL INPUT
RETURN returnType
LANGUAGE [c|java|python|sql]
AS '
    // source code
'.
```

Always called  
Null-check mandatory in code

# How to create an UDF ?

If any input is null, code block is skipped and return null

# How to create an UDF ?

```
CREATE [OR REPLACE] FUNCTION [IF NOT EXISTS]
[keyspace.]functionName (param1 type1, param2 type2, ...)
CALLED ON NULL INPUT | RETURNS NULL ON NULL INPUT
RETURN returnType
LANGUAGE language
AS '
// source code
';
```

## CQL types

- primitives (boolean, int, ...)
- collections (list, set, map)
- tuples
- UDT

# How to create an UDF ?

```
CREATE [OR REPLACE] FUNCTION [IF NOT EXISTS]
[keyspace.]functionName (param1 type1, param2 type2, ...)
CALLED ON NULL INPUT | RETURNS NULL ON NULL INPUT
RETURN returnType
LANGUAGE language
AS '
// source code here
'.
.'
```

## JVM supported languages

- Java, Scala
- Javascript (slow)
- Groovy, Jython, JRuby
- Clojure ( JSR 223 impl issue)

# How to create an UDF ?

```
CREATE [OR REPLACE] FUNCTION [IF NOT EXISTS]
[keyspace.]functionName (param1 type1, param2 type2, ...)
CALLED ON NULL INPUT | RETURNS NULL ON NULL INPUT
RETURN returnType
LANGUAGE language
AS '
// source code here
'.
```

Avoid simple quotes ( ' ) or escape them



# Simple UDF Demo

## $\max()$

# UDF creation for C\* 2.2 and 3.0

- UDF definition sent to **coordinator**
- Check for user permissions (**ALTER**, **DROP**, **EXECUTE**)
- Coordinator compiles it using **ECJ compiler** (not Javac because of **licence issue**).
- If language != Java, use other compilers (you should ship the jars on C\* servers)
- Announces SchemaChange = **CREATED**, Target = **FUNCTION** to **all nodes**
- Code loaded into **child classloader** of current classloader

# UDF creation in C\* 3.0

- Compiled byte code is verified
  - **static** and instance code block
  - declaring methods/constructor → no anonymous class creation
  - declaring inner class
  - use of **synchronized**
  - accessing forbidden classes, ex: java.net.NetworkInterface
  - accessing forbidden methods, ex: Class.forName()
  - ...
- Code loaded into **seperated classloader**
  - verify loaded classes against **whitelist/backlist**

# UDF execution in C\* 2.2

- No safeguard, arbitrary code execution server side ...
- Flag “`enable_user_defined_functions`” in `cassandra.yaml` turned **OFF** by default
- Keep your UDFs **pure (stateless)**
- Avoid expensive computation (for loop ...)

# UDF execution in C\* 3.0

- If "enable\_user\_defined\_functions\_threads" turned ON (default)
  - execute UDF in separated ThreadPool
  - if duration longer than "user\_defined\_function\_warn\_timeout", retry with new timeout = fail\_timeout – warn\_timeout
  - if duration longer than "user\_defined\_function\_fail\_timeout", exception
- Else
  - execute in same thread

# UDF creation best practices

- Keep your UDFs **pure (stateless)**
- No **expensive computation** (huge for loop)
- Prefer **Java** language for perf

# UDA

- Real use-case for UDF
- Aggregation server-side → huge network bandwidth saving
- Provide similar behavior for Group By, Sum, Avg etc ...

# How to create an UDA ?

```
CREATE [OR REPLACE] AGGREGATE [IF NOT EXISTS]  
[keyspace.]aggregateName(type1, type2, ...)  
SFUNC accumulatorFunction  
STYPE stateType  
[FINALFUNC finalFunction]  
INITCOND initCond;
```

Initial state type

Only type, no param name

State type

# How to create an UDA ?

```
CREATE [OR REPLACE] AGGREGATE [IF NOT EXISTS]  
[keyspace.]aggregateName(type1, type2, ...)  
SFUNC accumulatorFunction  
STYPE stateType  
[FINALFUNC finalFunction]  
INITCOND initCond;
```

Accumulator function. Signature:  
accumulatorFunction(stateType, type<sub>1</sub>, type<sub>2</sub>, ...)  
RETURNS stateType

# How to create an UDA ?

```
CREATE [OR REPLACE] AGGREGATE [IF NOT EXISTS]
```

```
[keyspace.]aggregateName(type1, type2, ...)
```

```
SFUNC accumulatorFunction
```

```
STYPE stateType
```

```
[FINALFUNC finalFunction]
```

```
INITCOND initCond;
```

Optional final function. Signature:  
finalFunction(stateType)

# How to create an UDA ?

```
CREATE [OR REPLACE] AGGREGATE [IF NOT EXISTS]
[keyspace.]aggregateName(type1, type2, ...)
SFUNC accumulatorFunction
STYPE stateType
[FINALFUNC finalFunction]
INITCOND initCond;
```

UDA return type ?  
If finalFunction  
• return type of finalFunction  
Else  
• return stateType

# How to create an UDA ?

```
CREATE [OR REPLACE] AGGREGATE [IF NOT EXISTS]
[keyspace.]aggregateName(type1, type2, ...)
SFUNC accumulatorFunction
STYPE stateType
[FINALFUNC finalFunction]
INITCOND initCond;
```

- If accumulatorFunction RETURNS NULL ON NULL INPUT
  - initCond should be defined
- Else
  - initCond can be null (default value)

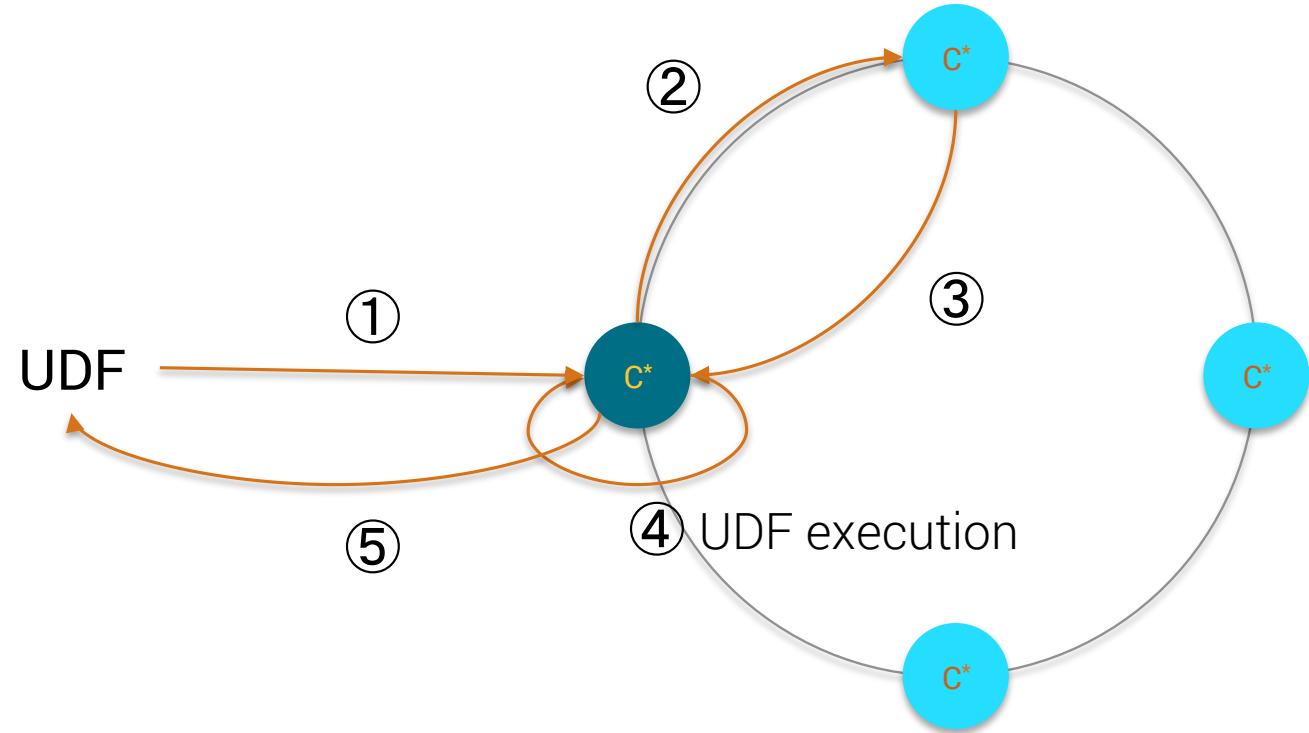
# UDA execution

- Start with **state** = **initCond**
- For each row
  - call accumulator function with **previous state + params**
  - update **state** with accumulator function output
- If no **finalFunction**
  - return **last state**
- Else
  - apply **finalFunction** and returns its output

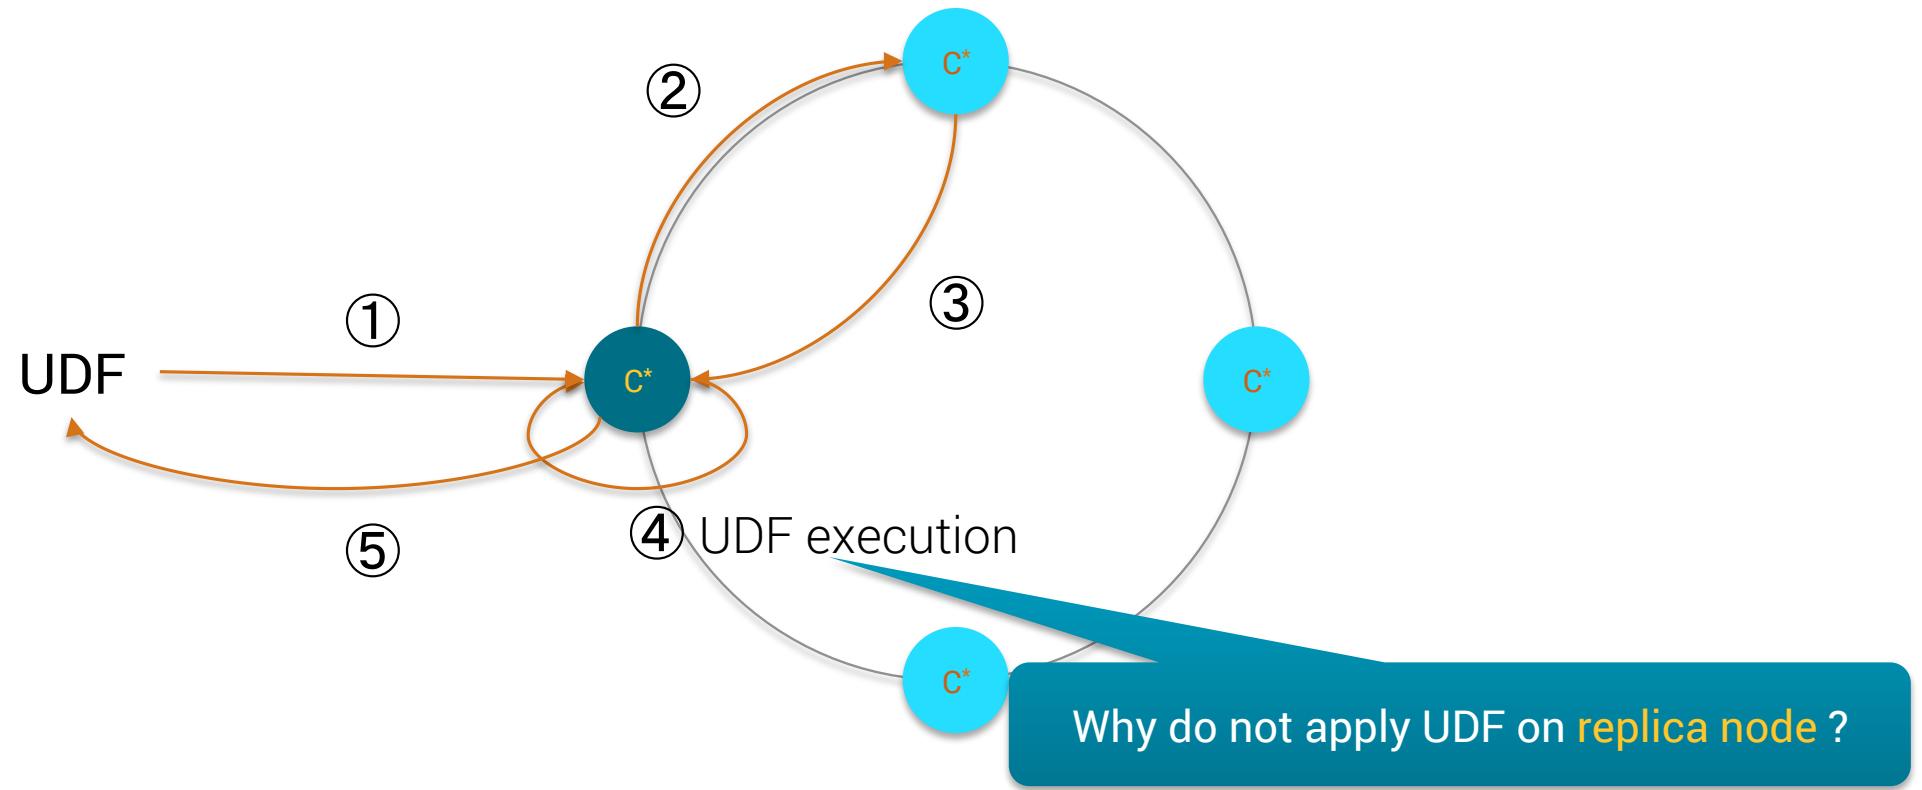


# UDA Demo Group By

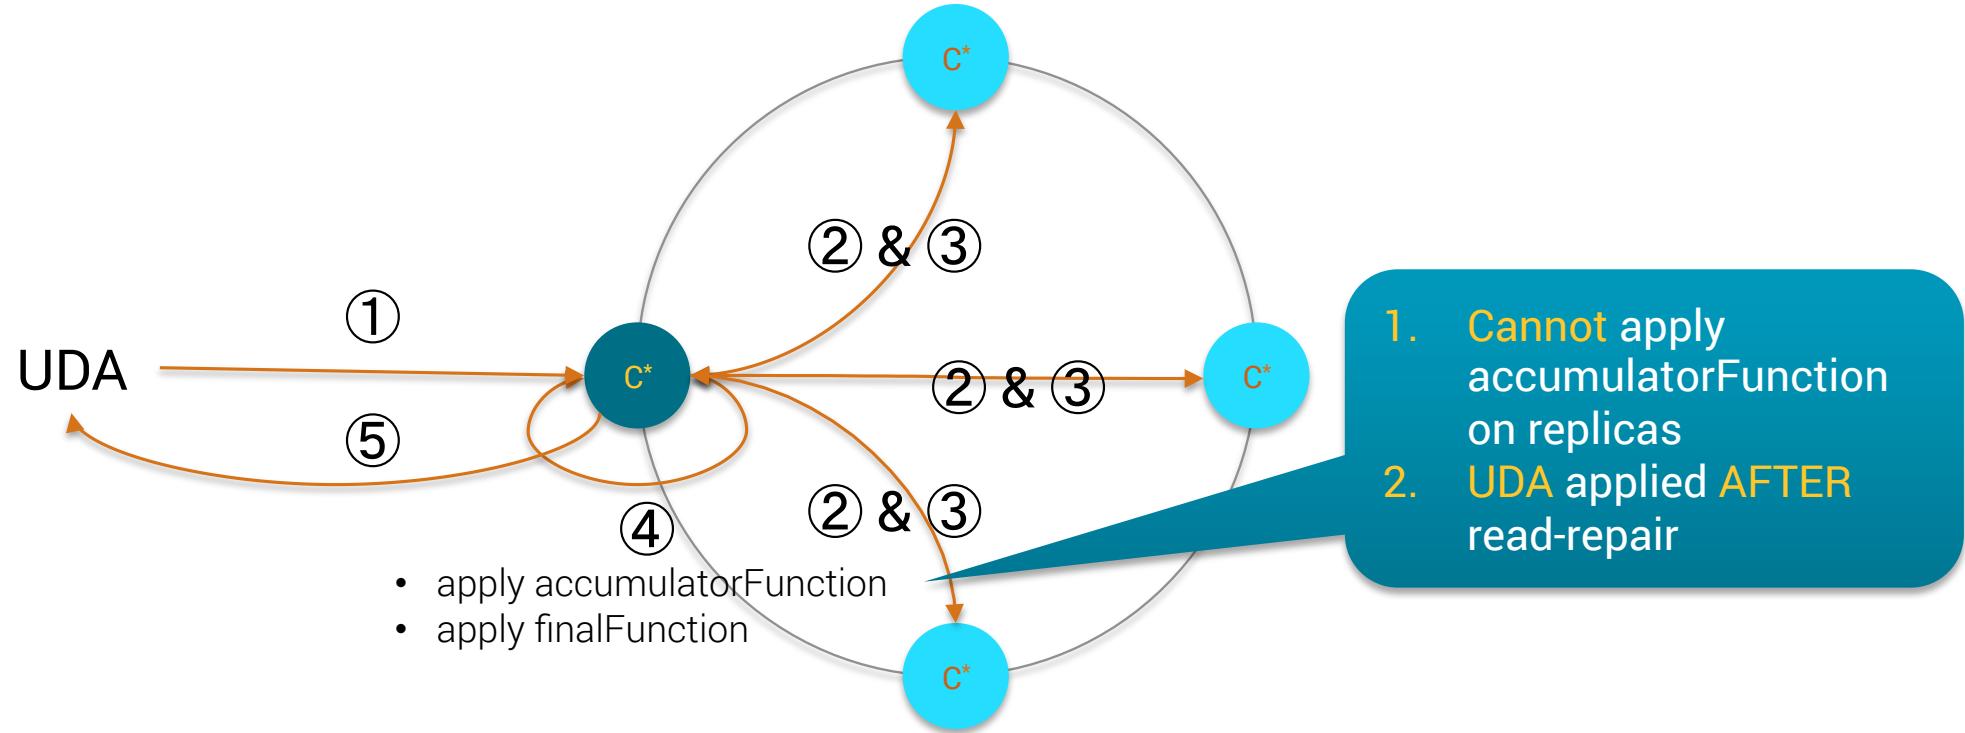
# Gotchas



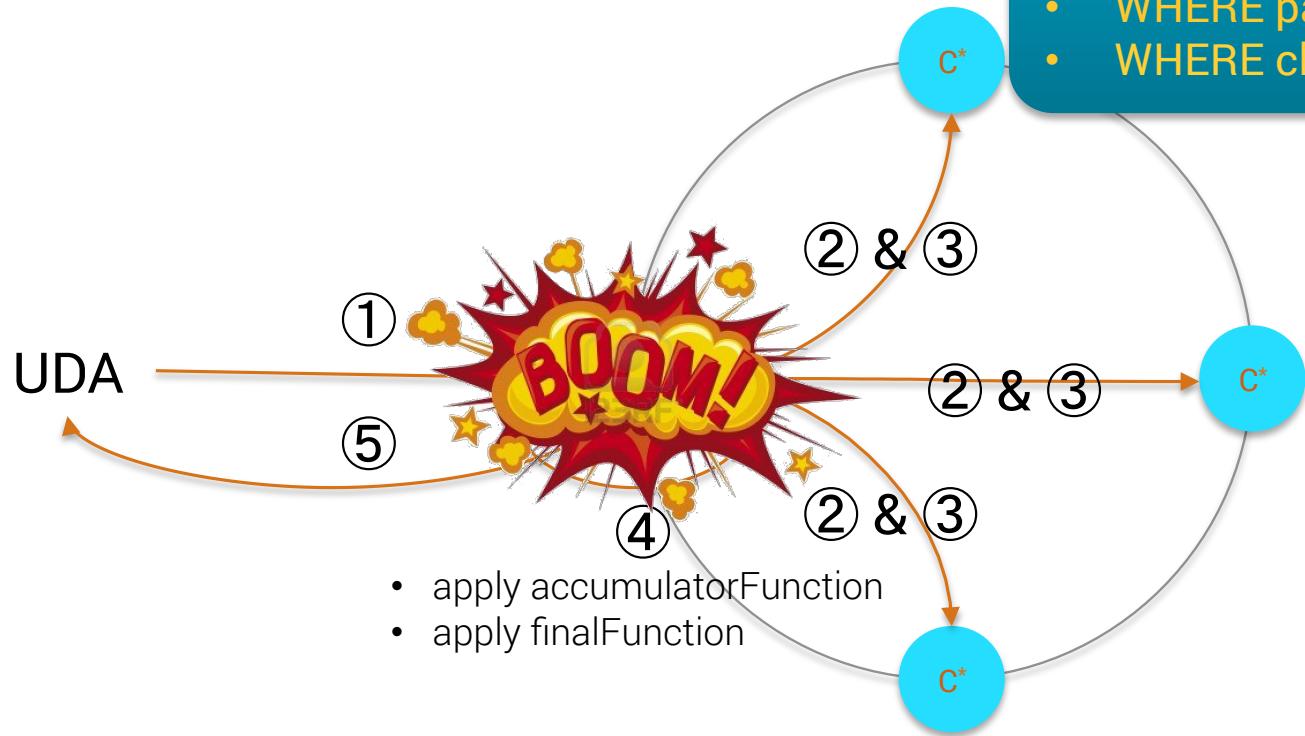
# Gotchas



# Gotchas



# Gotchas

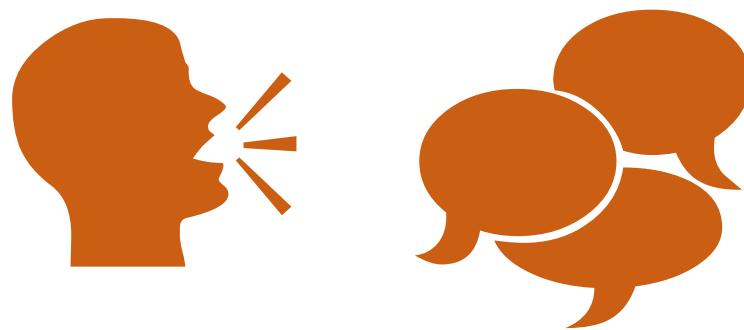


Avoid UDA on large number of rows

- WHERE partitionKey IN (...)
- WHERE clustering >= AND clustering <=

# Future applications of UDFs & UDAs

- Functional index (**CASSANDRA-7458**)
- Partial index (**CASSANDRA-7391**)
- WHERE clause filtering with UDF
- ...



# Q & A



# Materialized Views

- Why ?
- Detailed Impl
- Gotchas

# Why Materialized Views ?

- Relieve the pain of manual denormalization

```
CREATE TABLE user(  
    id int PRIMARY KEY,  
    country text,  
    ...  
);  
CREATE TABLE user_by_country(  
    country text,  
    id int,  
    ...,  
    PRIMARY KEY(country, id)  
);
```

# Materialized View In Action

```
CREATE MATERIALIZED VIEW user_by_country  
AS SELECT country, id, firstname, lastname  
FROM user  
WHERE country IS NOT NULL AND id IS NOT NULL  
PRIMARY KEY(country, id)
```



```
CREATE TABLE user_by_country (  
    country text,  
    id int,  
    firstname text,  
    lastname text,  
    PRIMARY KEY(country, id));
```

# Materialized View Syntax

CREATE MATERIALIZED VIEW [IF NOT EXISTS]

keyspace\_name.view\_name

Must select **all primary key columns** of base table

AS SELECT column<sub>1</sub>, column<sub>2</sub>, ...

FROM keyspace\_name.table\_name

WHERE column<sub>1</sub> IS NOT NULL AND column<sub>2</sub> IS NOT NULL ...

PRIMARY KEY(column<sub>1</sub>, column<sub>2</sub>, ...)

- **IS NOT NULL** condition for now
- more complex conditions in future

- at least **all primary key columns** of base table (ordering can be different)
- maximum 1 column NOT pk from base table
- 1:1 relationship between base row & MV row

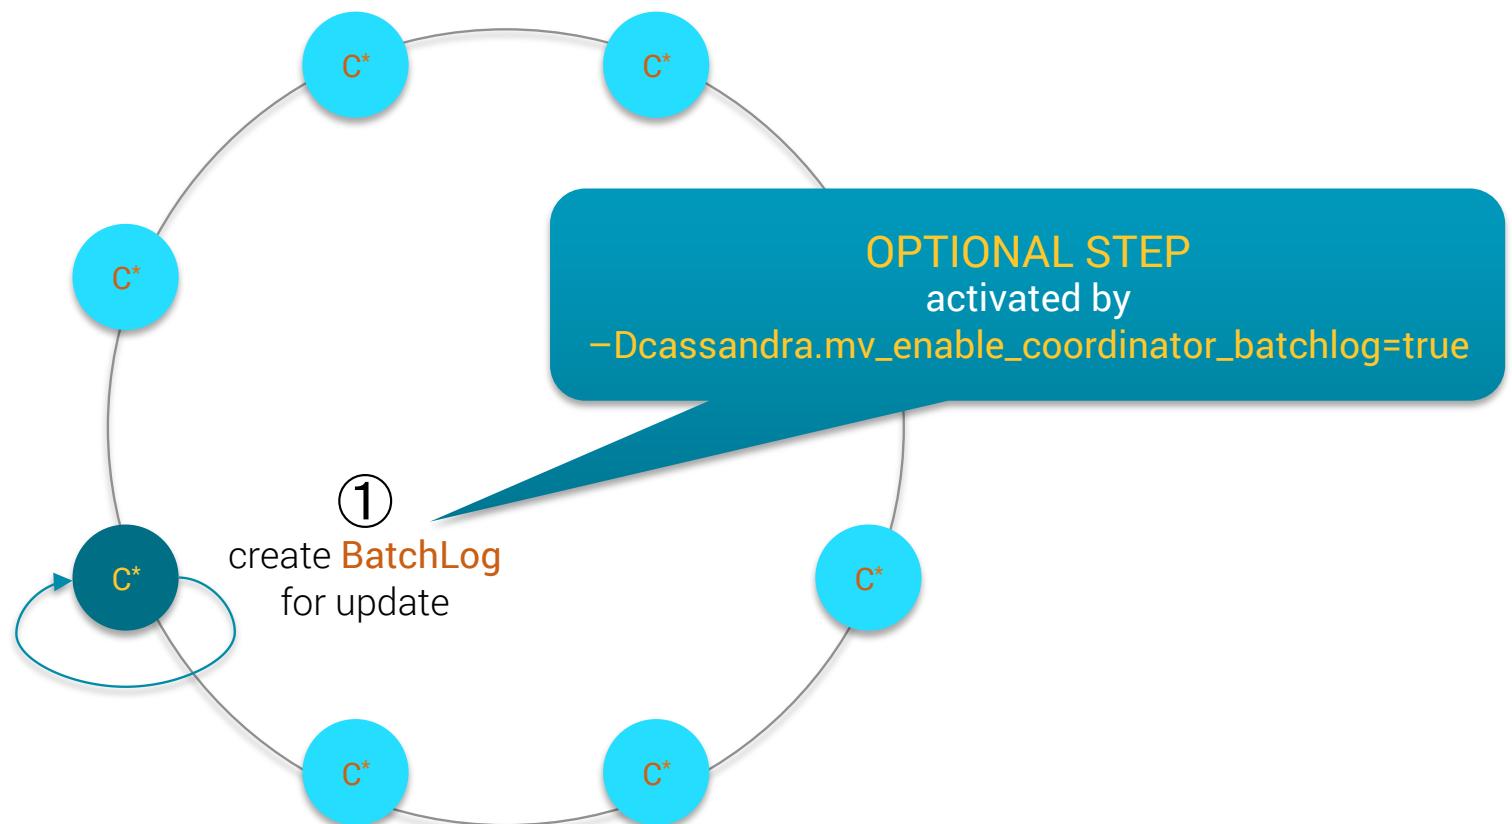


# Materialized View Demo

## user\_by\_country

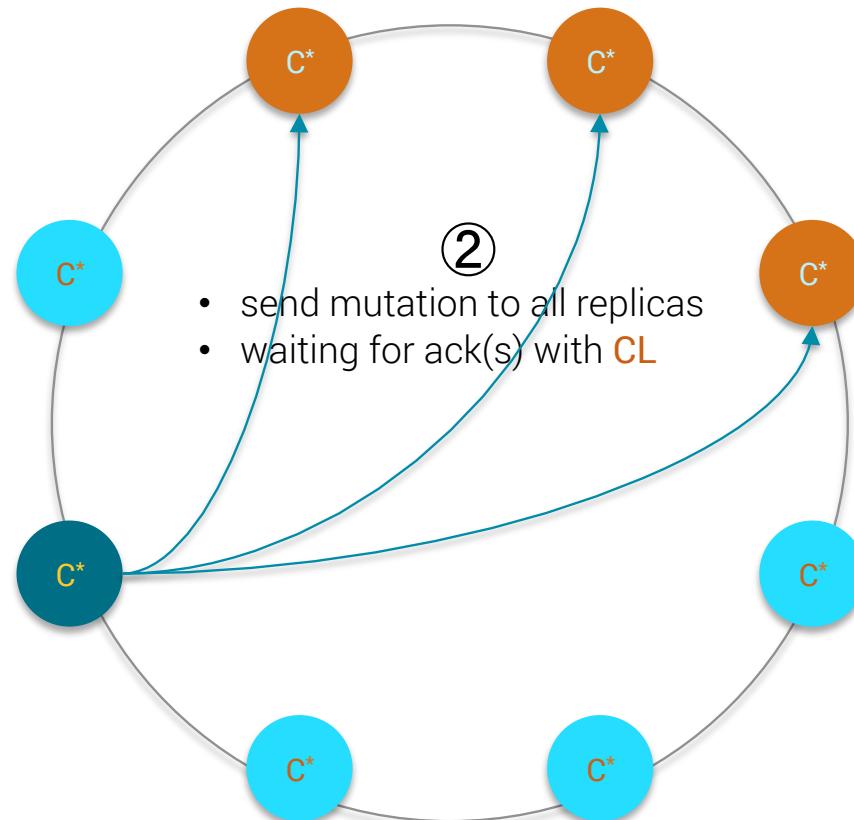
# Materialized View Impl

```
UPDATE user  
SET country='FR'  
WHERE id=1
```



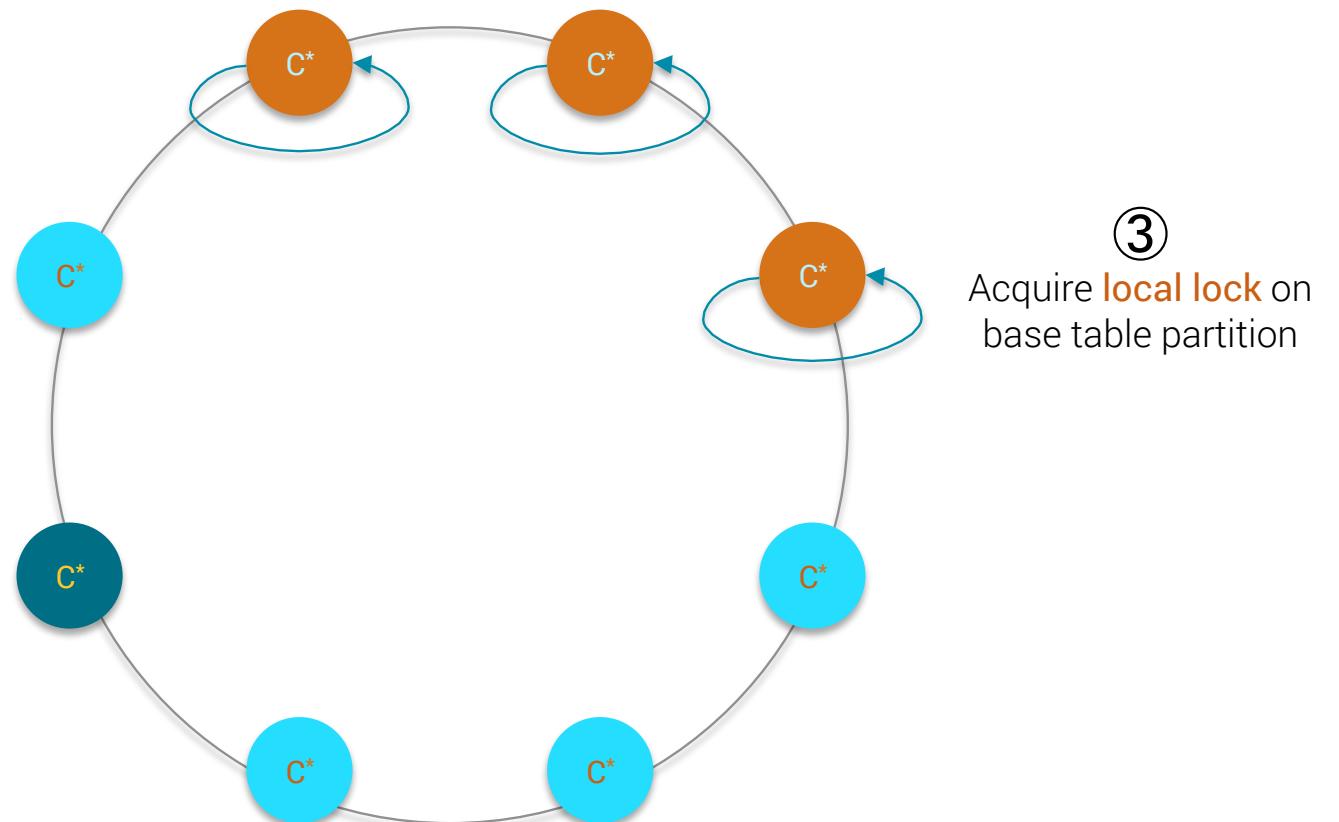
# Materialized View Impl

```
UPDATE user  
SET country='FR'  
WHERE id=1
```



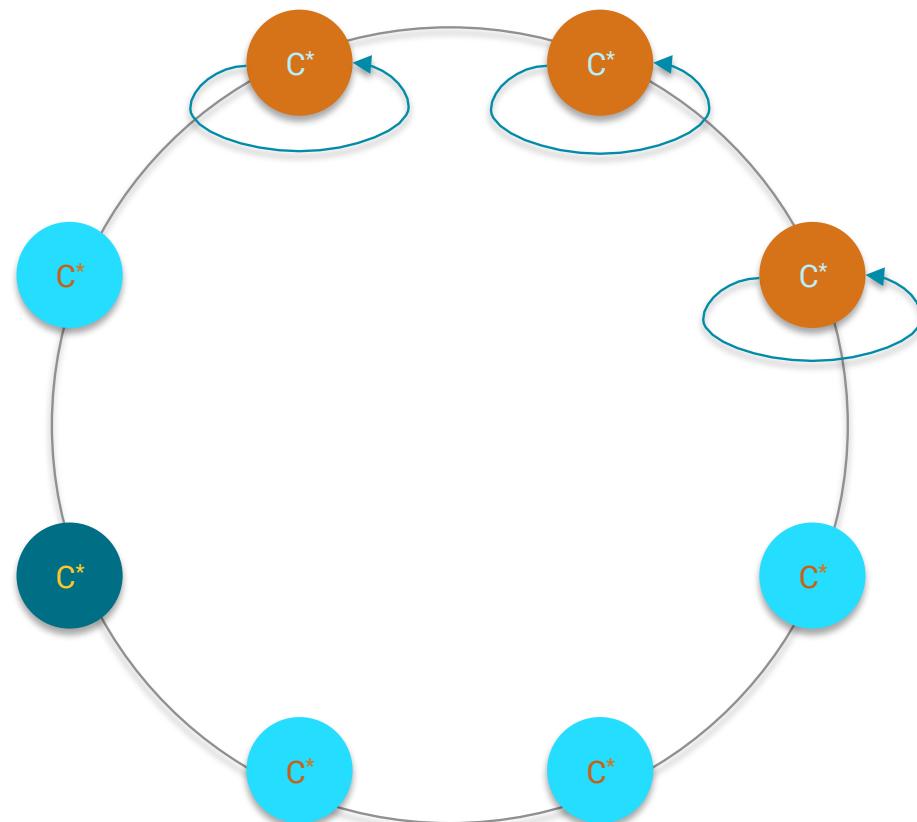
# Materialized View Impl

```
UPDATE user  
SET country='FR'  
WHERE id=1
```



# Materialized View Impl

```
UPDATE user  
SET country='FR'  
WHERE id=1
```

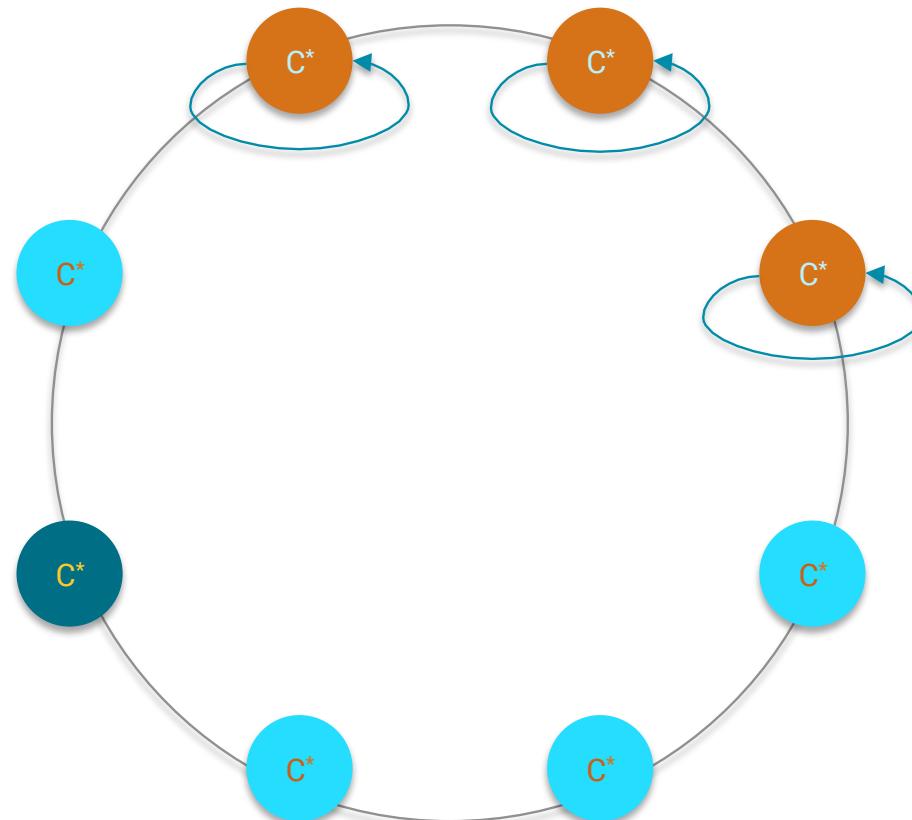


④

Local read to fetch current values  
**SELECT \* FROM user WHERE id=1**

# Materialized View Impl

UPDATE user  
SET country='FR'  
WHERE id=1

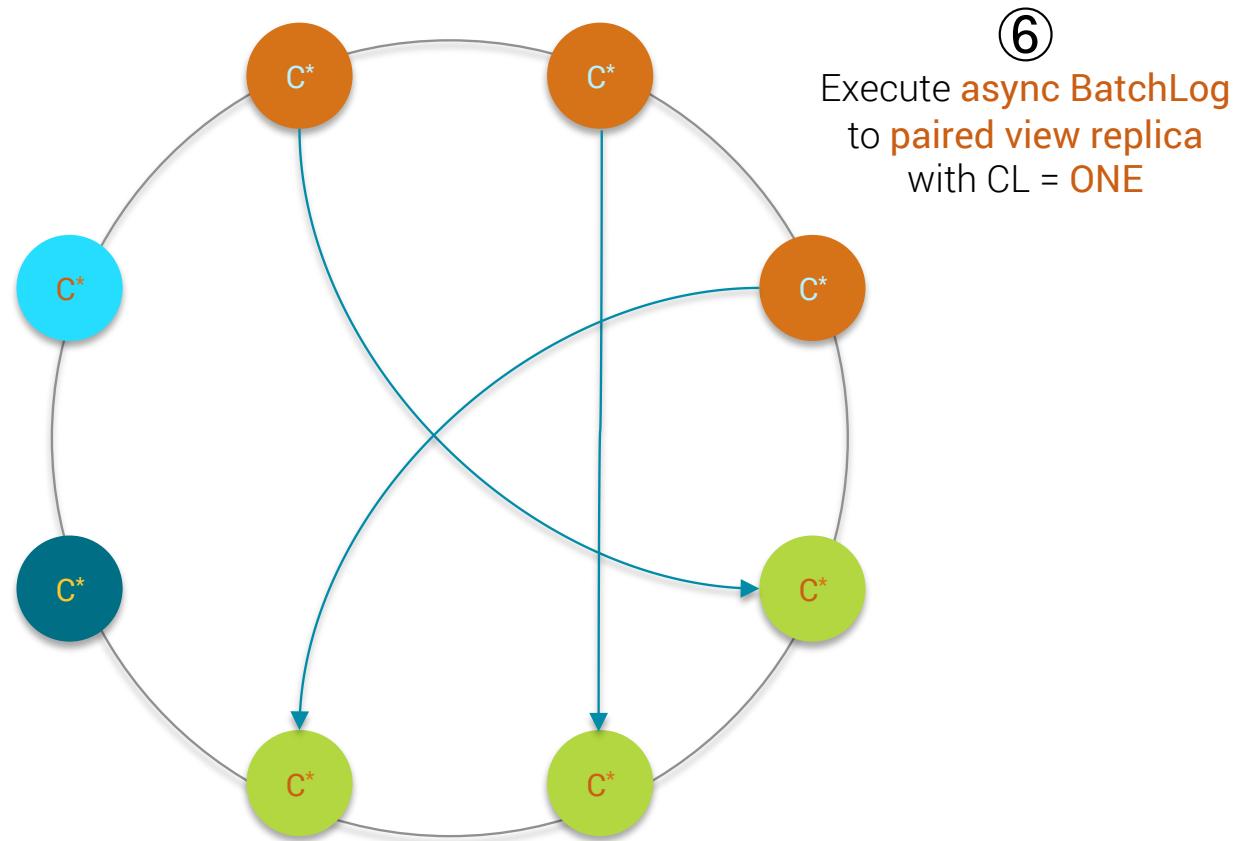


⑤ Create BatchLog with

- `DELETE FROM user_by_country WHERE country = 'old_value'`
- `INSERT INTO user_by_country(country, id, ...)`  
`VALUES('FR', 1, ...)`

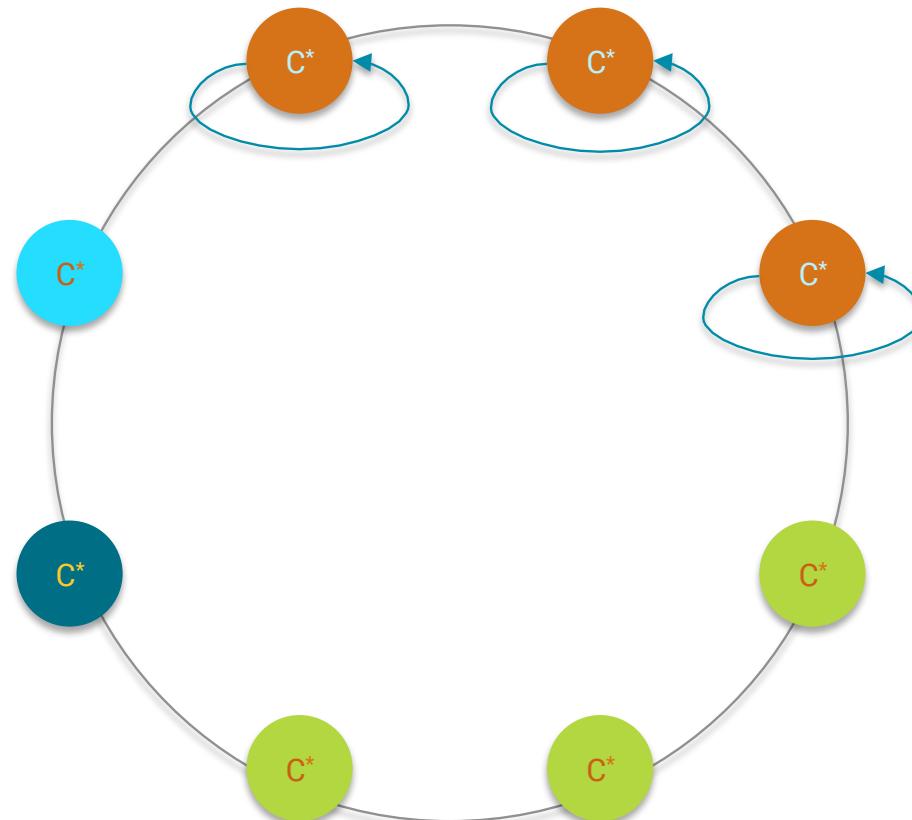
# Materialized View Impl

UPDATE user  
SET country='FR'  
WHERE id=1



# Materialized View Impl

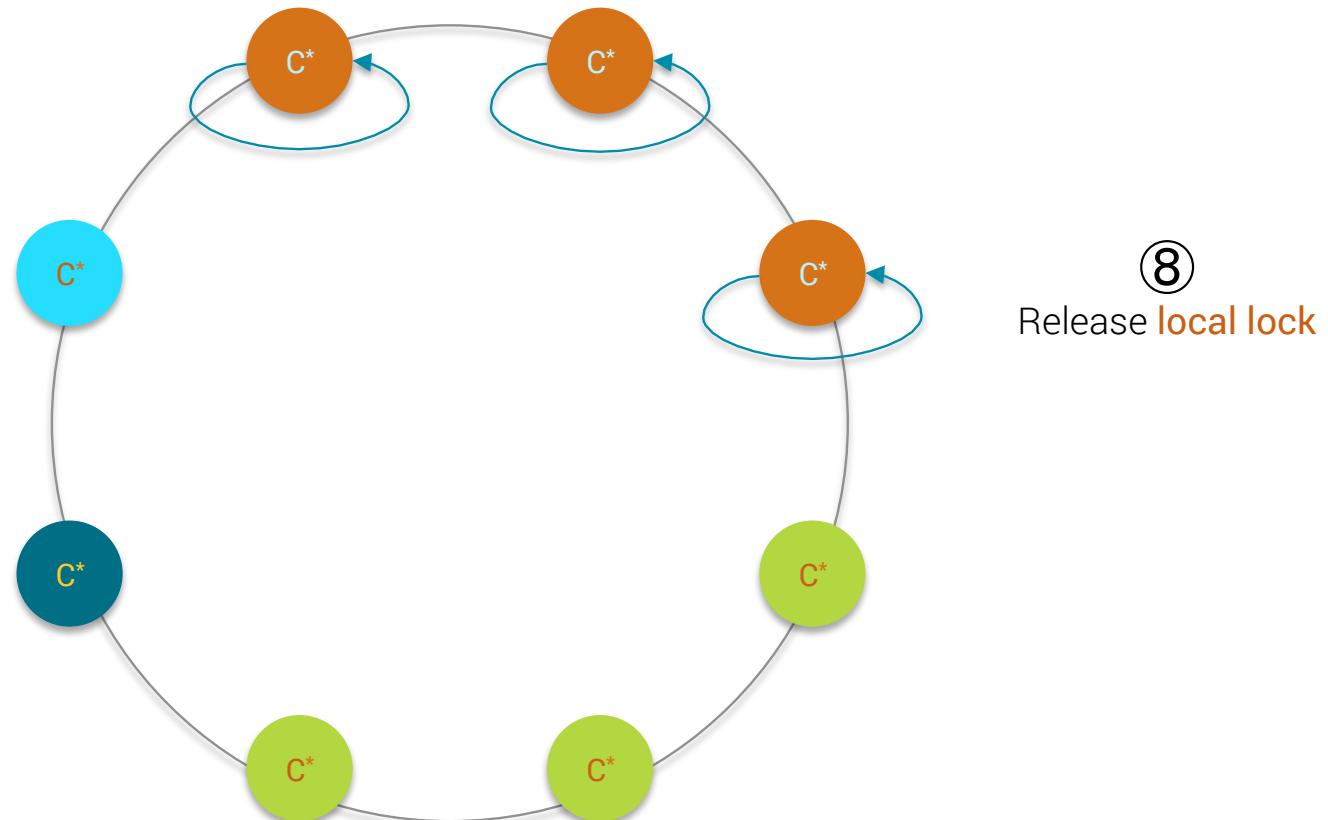
UPDATE user  
SET country='FR'  
WHERE id=1



7  
Apply base table update locally  
SET COUNTRY='FR'

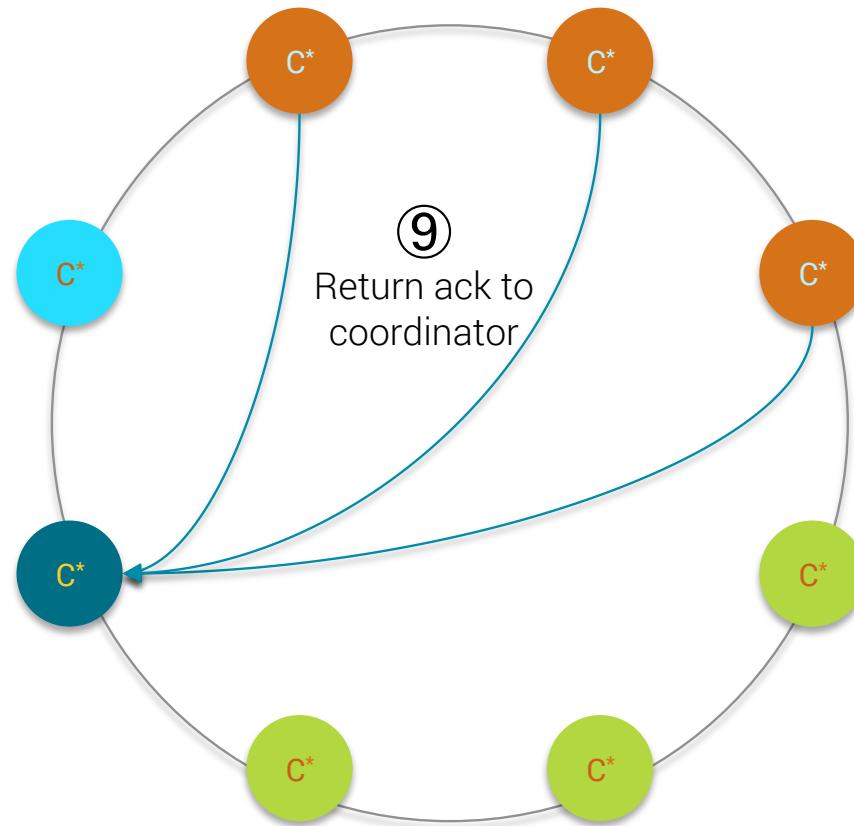
# Materialized View Impl

UPDATE user  
SET country='FR'  
WHERE id=1



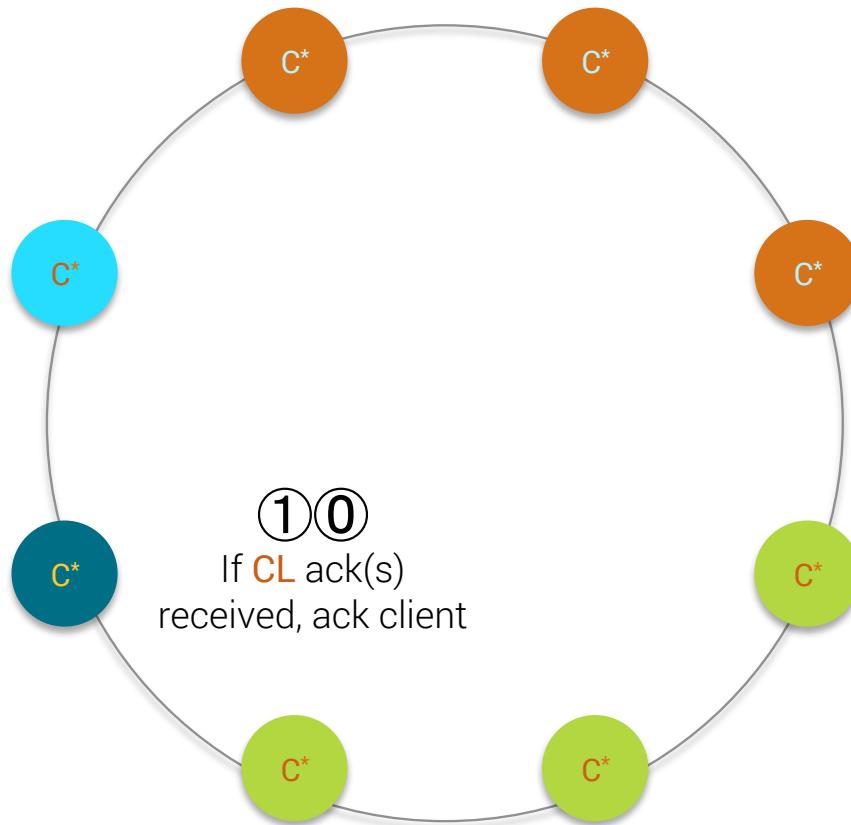
# Materialized View Impl

```
UPDATE user  
SET country='FR'  
WHERE id=1
```



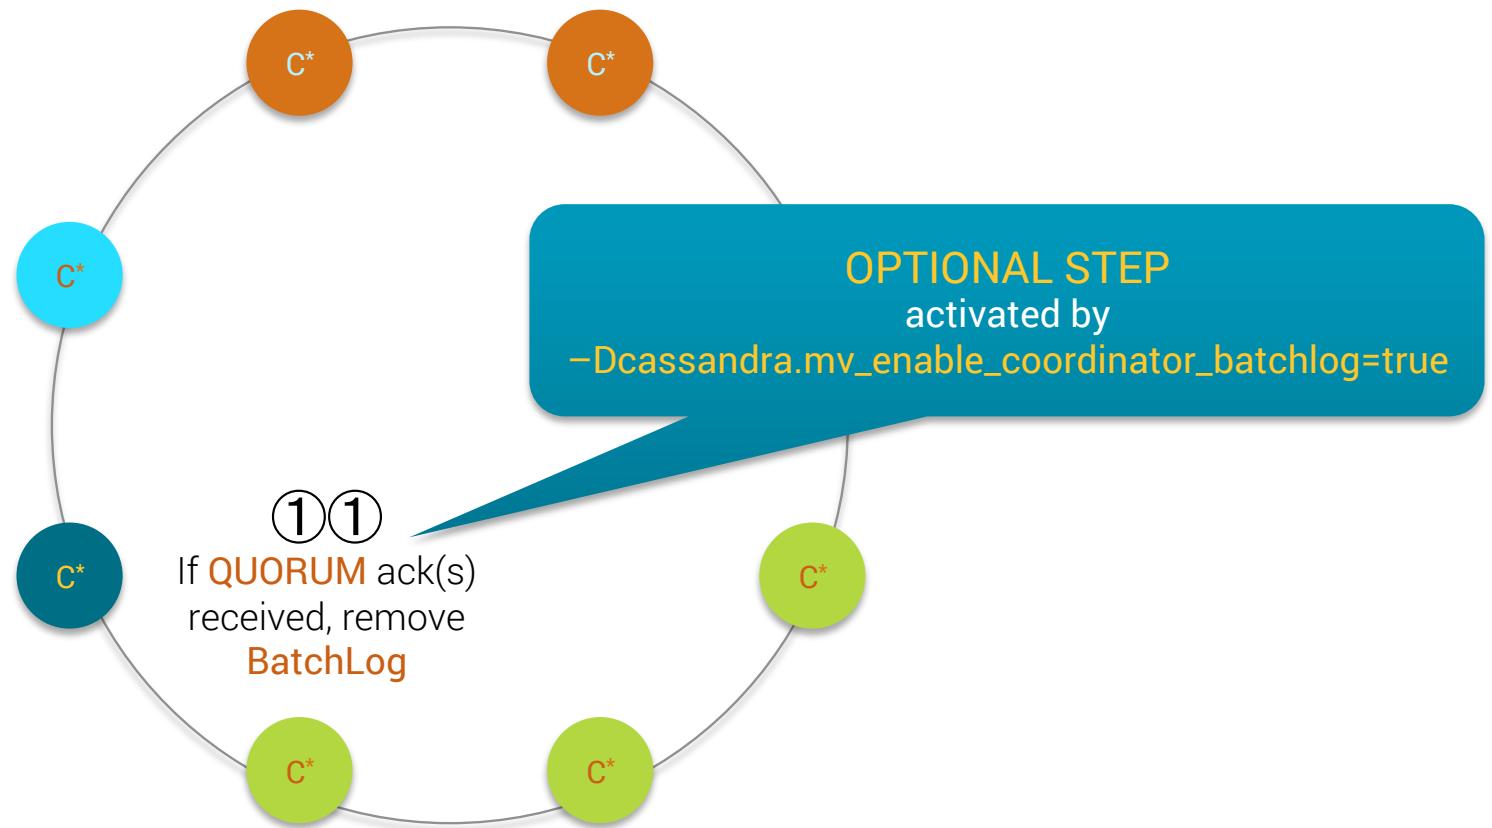
# Materialized View Impl

UPDATE user  
SET country='FR'  
WHERE id=1



# Materialized View Impl

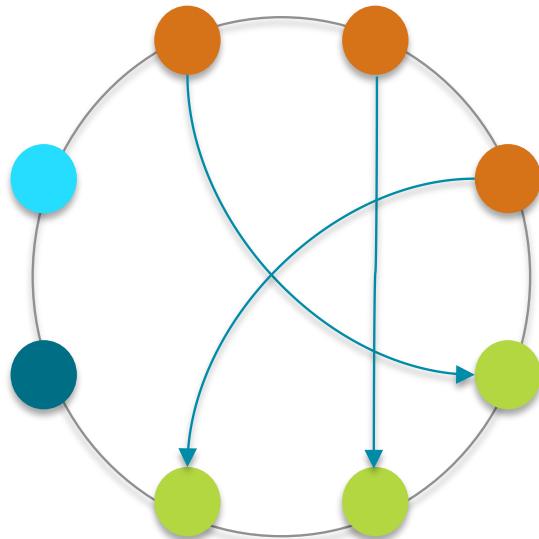
UPDATE user  
SET country='FR'  
WHERE id=1



# Materialized Views impl explained

What is paired replica ?

- Base primary replica for **id=1** is paired with MV primary replica for **country='FR'**
- Base secondary replica for **id=1** is paired with MV secondary replica for **country='FR'**
- etc ...



# Materialized Views impl explained

Why **local lock** on base replica ?

- to provide **serializability** on concurrent updates

Why **BatchLog** on base replica ?

- necessary for **eventual durability**
- replay the MV delete + insert until successful

Why **BatchLog** on base replica uses **CL = ONE** ?

- each base replica is responsible for update of its **paired** MV replica
- CL > ONE will create un-necessary **duplicated mutations**

# Materialized Views impl explained

Why **BatchLog** on coordinator with **QUORUM** ?

- necessary to cover some **edge-cases**

# MV Failure Cases: concurrent updates

*Without Local Lock*

## 1) UPDATE ... SET country='US'

Read base row (country='UK')  
• DELETE FROM mv WHERE country='UK'  
• INSERT INTO mv ... (country)  
VALUES('US')  
• Send async **BatchLog**  
• Apply update country='US'

## 2) UPDATE ... SET country='FR'

Read base row (country='UK')  
• DELETE FROM mv WHERE country='UK'  
• INSERT INTO mv ... (country)  
VALUES('FR')  
• Send async **BatchLog**  
• Apply update country='FR'



# MV Failure Cases: concurrent updates *with local lock*

## 1) UPDATE ... SET country='US'



- Read base row (country='UK')
- DELETE FROM mv WHERE country='UK'
- INSERT INTO mv ... (country) VALUES('US')
- Send async **BatchLog**
- Apply update country='US'



## 2) UPDATE ... SET country='FR'

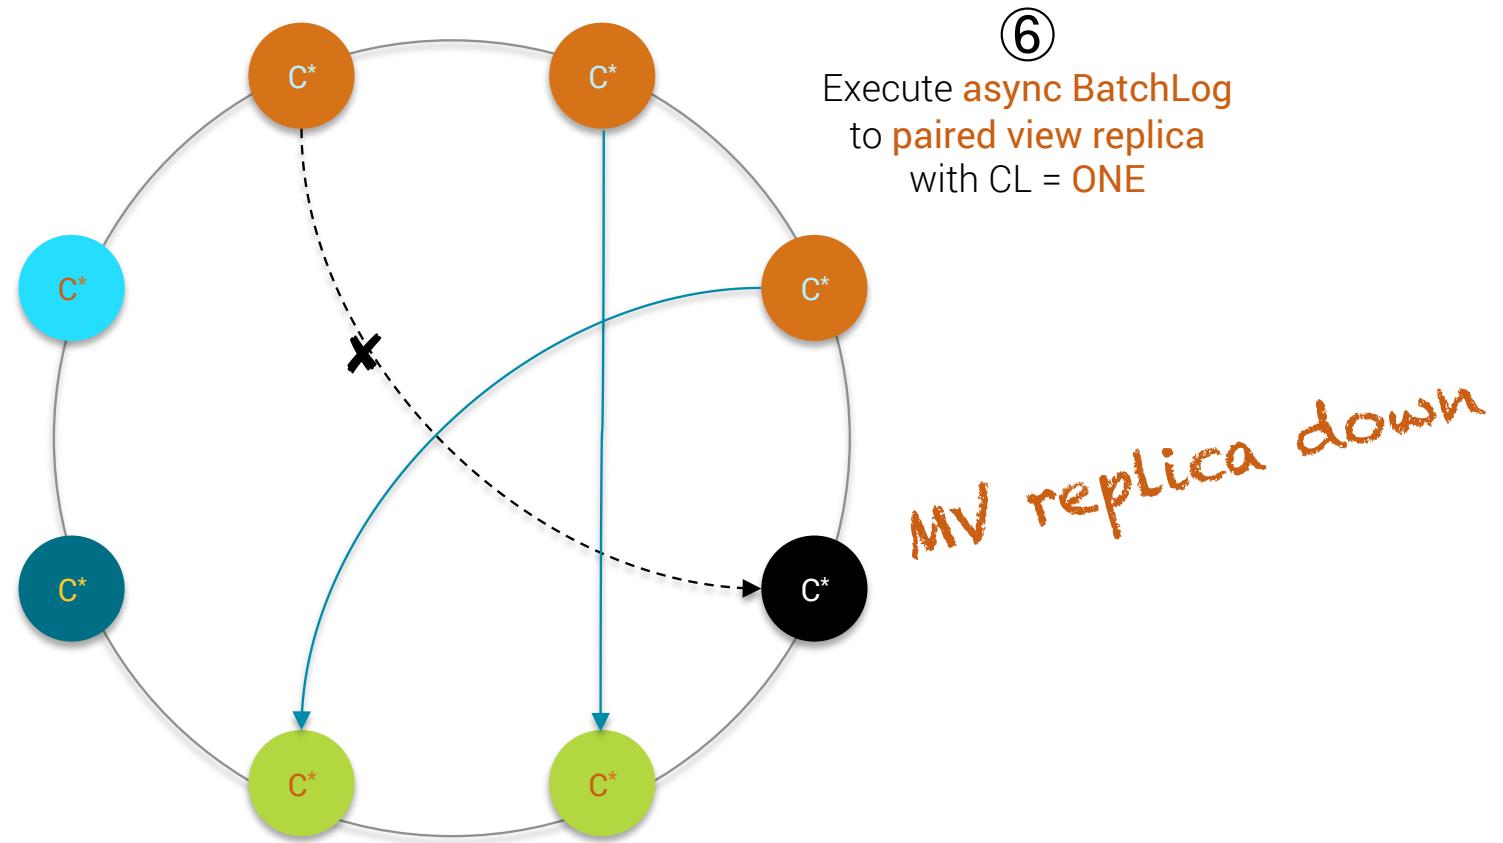


- Read base row (country='US')
- DELETE FROM mv WHERE country='US'
- INSERT INTO mv ... (country) VALUES('FR')
- Send async **BatchLog**
- Apply update country='FR'

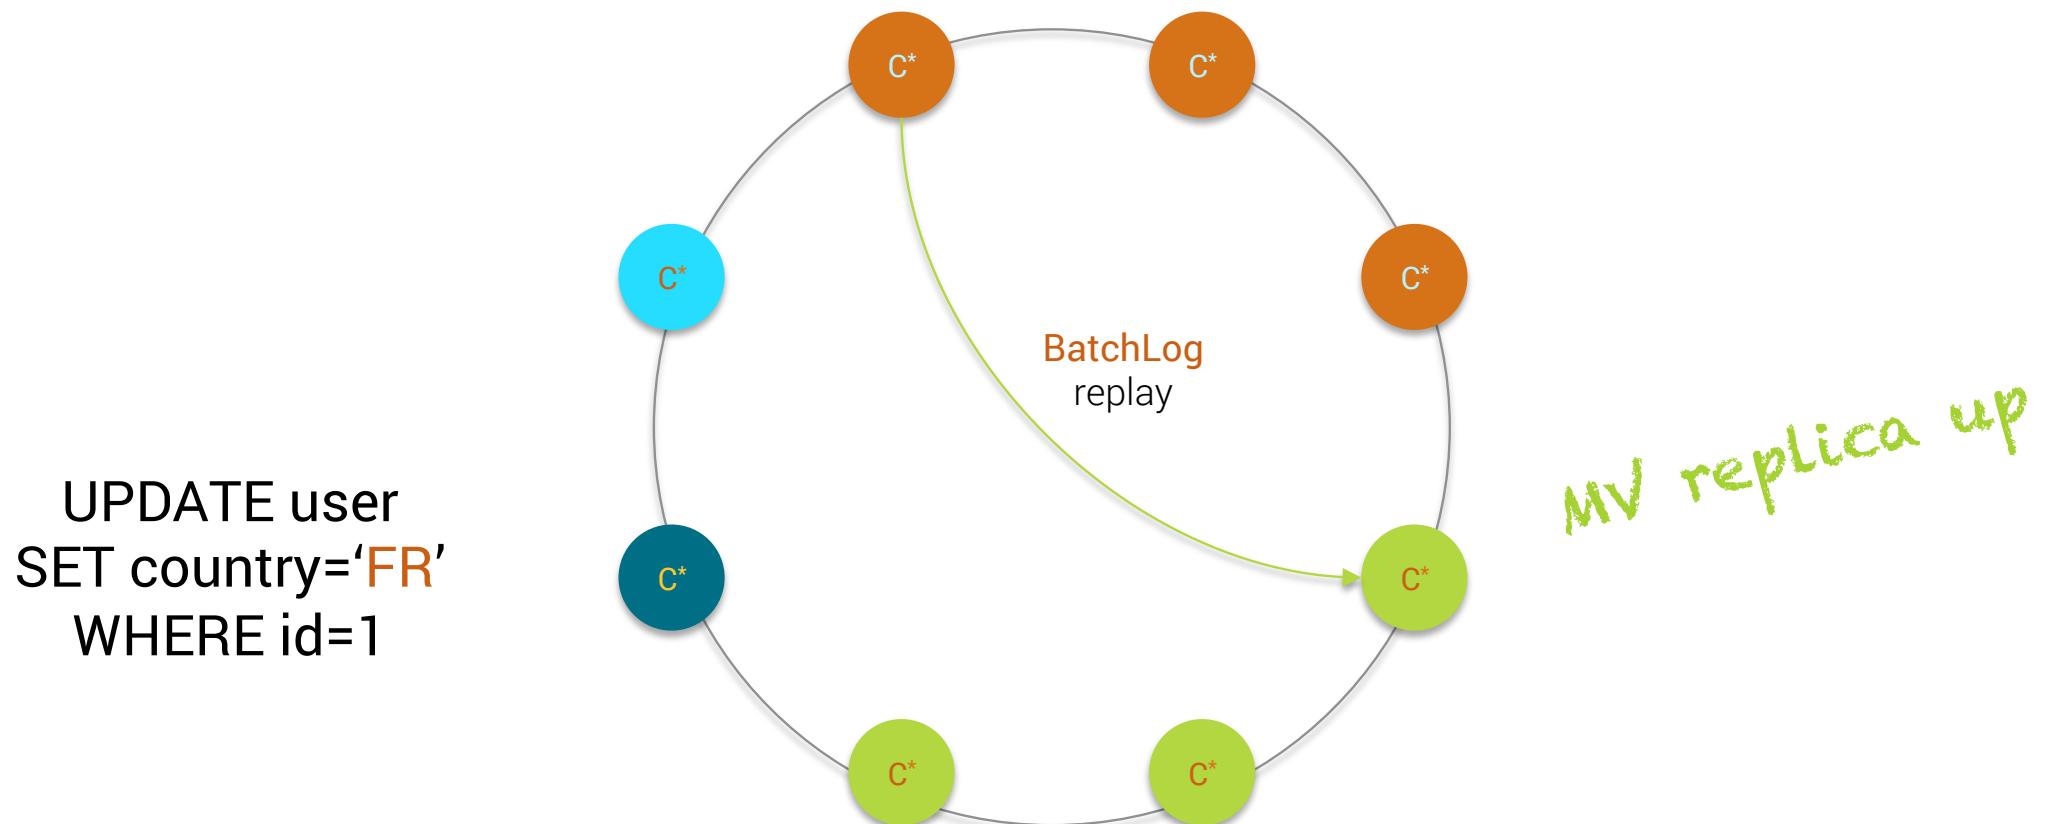


# MV Failure Cases: failed updates to MV

```
UPDATE user  
SET country='FR'  
WHERE id=1
```



# MV Failure Cases: failed updates to MV



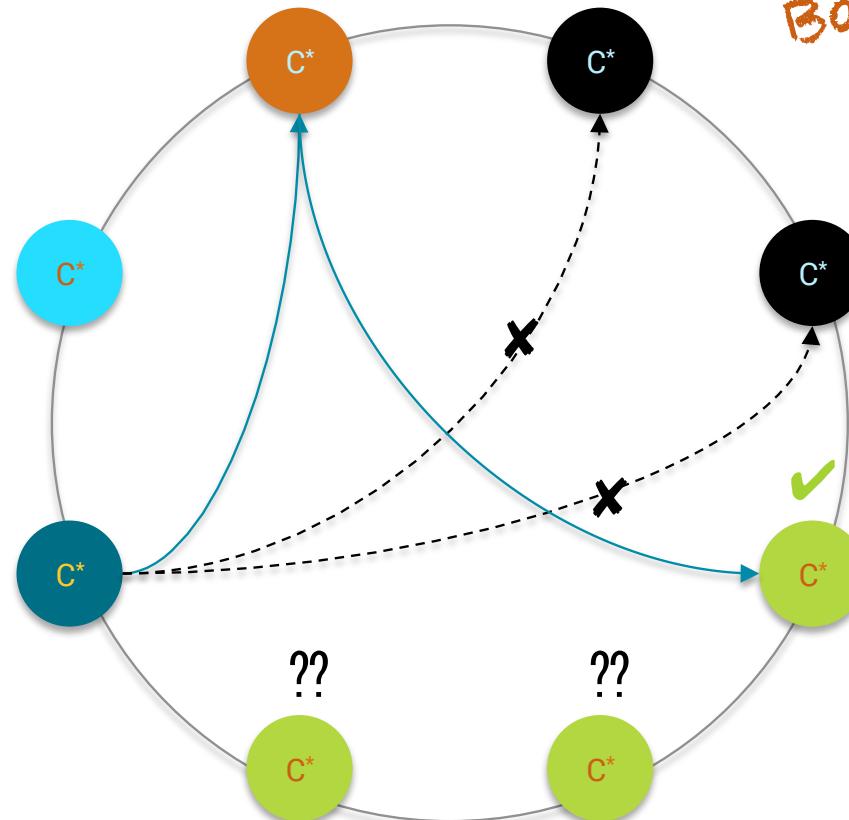
# MV Failure Cases: base replica(s) down

Base replicas down

②

- send mutation to all replicas
- waiting for ack(s) with **CL**

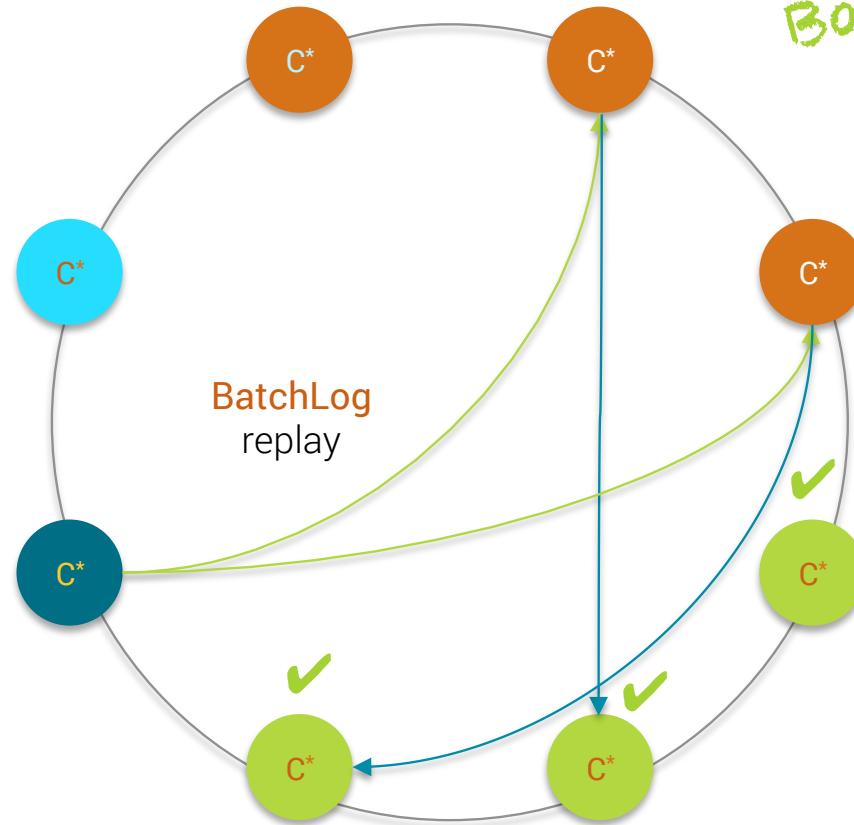
UPDATE user  
SET country='FR'  
WHERE id=1



# MV Failure Cases: base replica(s) down

Base replicas up

UPDATE user  
SET country='FR'  
WHERE id=1



# Materialized Views Gotchas

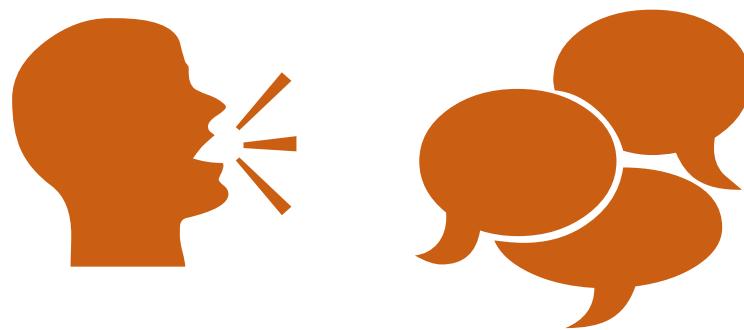
- Write performance
  - local lock
  - local read-before-write for MV (*most of perf hits here*)
  - local batchlog for MV
  - coordinator batchlog (**OPTIONAL**)
  - ↗ you only pay this price **once** whatever number of MV
- Consistency level
  - CL honoured for base table, **ONE** for MV + local batchlog
  - **QUORUM** for coordinator batchlog (**OPTIONAL**)

# Materialized Views Gotchas

- Beware of hot spots !!!
  - MV `user_by_gender` 😱
- Repair, read-repair, index rebuild, decommission ...
  - repair on base replica → update on MV paired replica
  - repair on MV replica possible
  - read-repair on MV behaves as normal read-repair

# Materialized Views Gotchas

- Schema migration
  - cannot drop column from base table used by MV
  - can add column to base table, initial value = null from MV
  - cannot drop base table, drop all MVs first



# Q & A

# Thank You



@doanduyhai



duy\_hai.doan@datastax.com

<https://academy.datastax.com/>