



Cassandra for Developers

Module 2

Data Model

LXFT
LISTED
NYSE

Whole picture

- ◆ Column-oriented store
- ◆ Multidimensional map
- ◆ $\text{Map}<\text{RowKey}, \text{SortedMap}<\text{Name}, \text{Value}, \text{TimeStamp}, \text{TTL}>>$
- ◆ Efficient key lookup
- ◆ RowKey, Name, Value – byte arrays

Column

- ◆ Column is like struct

- Name

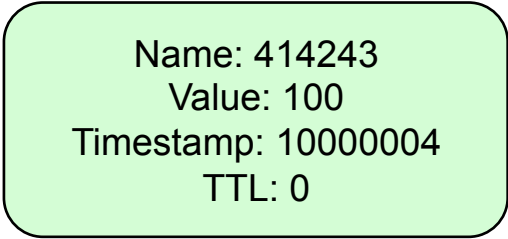
- ◆ Mandatory, any type and value, unique
 - ◆ Stored with every column!
 - ◆ Data type is called a comparator
 - ◆ Max 64Kb

- Value

- ◆ Any type and value (effectively byte array)
 - ◆ Can be empty
 - ◆ Data type is called a validator
 - ◆ Max 2 Gb

- Timestamp

- Time to live (TTL) – optional (def: 0)

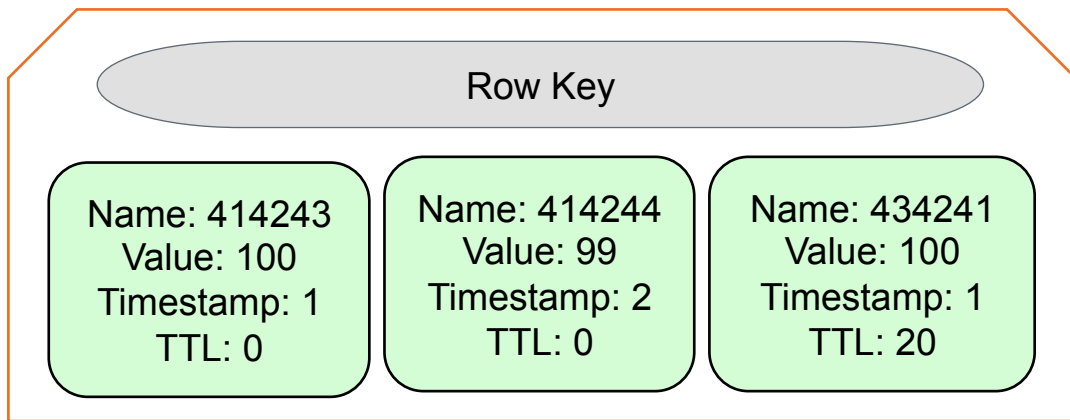


Name: 414243
Value: 100
Timestamp: 10000004
TTL: 0

Rows

- ◆ Row is like named hash table

- Name = row key
 - ◆ Unique in a column family
- Collection of columns (indexed by name)
 - ◆ Orderable




- ◆ Like RDBMS tuple

- ◆ Each row replica is stored on one node (no splitting)

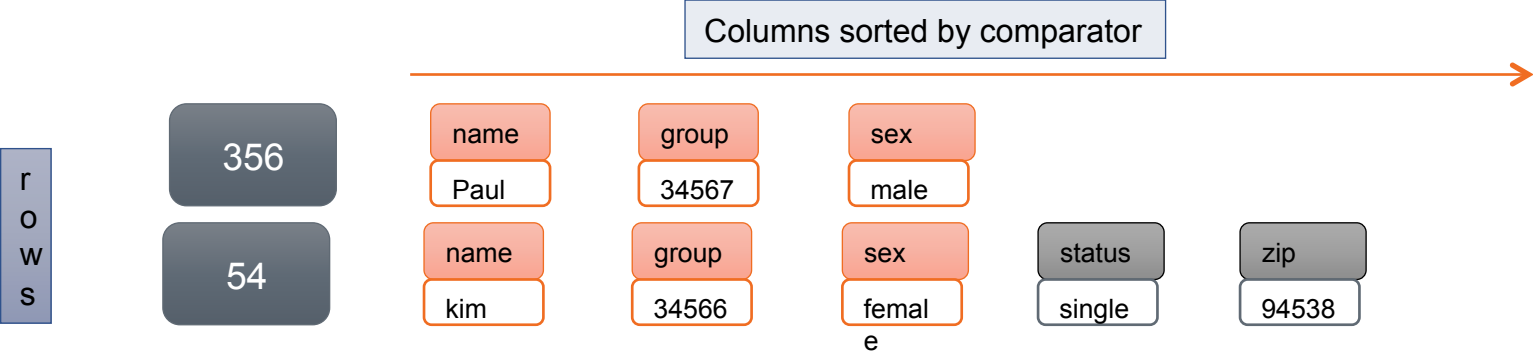
- Must fit on disk

Column Families


- ◆ Column family is like named hash table of hash tables
 - Name
 - Collection of rows (indexed by row keys)
 - ◆ Unique row keys
 - ◆ Columns of different rows may differ
 - i.e. CF is sparse nested hash table
 - ◆ Default ordering: by clustering columns – second part of row key
 - Compare row keys as bytes, long, UTF8, ASCII, UUID, custom
- ◆ Loose analog to RDBMS tables



Data Model



Composite columns



Static Column Family

- ◆ Most similar to RDBMS tables
- ◆ Most rows have the same column names
- ◆ But columns in rows can be different

Dynamic Column Family

- ◆ “Wide” rows
- ◆ Highly denormalized
- ◆ All columns needed to answer a query to a row

Super column

- ◆ Super column is like named map of columns
 - Name
 - Collection of columns (indexed by name)
- ◆ Deprecated since C* 2.0
- ◆ No access from CQL (only Thrift)
- ◆ Use composite columns instead

Super column family

- ◆ Like a hash table of hash tables of hash tables
- ◆ Super column family is a column family where rows consist of super columns instead of columns

Keyspace

- ◆ Keyspace is a collection of column families (and super column families)
- ◆ Has specific replication configuration
 - Replication strategy
 - Replication factor
- ◆ Like schema in Oracle Database
- ◆ Like database in MySQL, Microsoft SQL Server

Cluster

- ◆ Collection of keyspaces
- ◆ At least
 - system keyspace
 - user keyspaces

Partitioners

- ◆ Token = Partitioner(Key)
- ◆ Murmur3Partitioner (RandomPartitioner)
 - Ordering is random
- ◆ ByteOrderedPartitioner
 - Ordering is lexical for each node and the ring

Column families properties that drives data modeling

- ◆ Are not tables
- ◆ Can be very wide
- ◆ Can be very narrow
- ◆ There's no joins
 - While foreign keys are there
 - ◆ Value (or name, or name part) of a column might be an another row key
- ◆ There's only one index by default (row key)
 - Actually, used for distributing rows (=partitioning =sharding)

Secondary indexes

- ◆ Any column (or group of columns) can be indexed
- ◆ Performant for low-cardinality columns
- ◆ Local for node (not for keyspace or cluster)
- ◆ The only possibility to use a column in a search expression

Some basic modeling patterns

- ◆ Entities in column family
- ◆ (Another) column family as index
- ◆ Materialized views
- ◆ Time series
- ◆ Event sourcing

Entities modeling

- ◆ Single primary key
- ◆ Fairly consistent columns names
- ◆ Fairly narrow rows
- ◆ Feels like schema-less relational data model

Column family as index

- ◆ Rows are indexed by row key
- ◆ Secondary indexes prefer low cardinality
- ◆ We have to build our own
- ◆ Query-first data modeling
 - Column family per query
 - Column values become entity row keys

Materialized views

- ◆ Based on entity data
- ◆ A separate column family
- ◆ Entity data organized by indexed key

Event sourcing

- ◆ Martin Fowler's Enterprise Application Architecture pattern: "Capture all changes to an application state as a sequence of events"
- ◆ Persist state change events
- ◆ Do not persist current state
- ◆ All writes are immutable
- ◆ At any time events sequence can be replayed and current state will be formed

Time series

- ◆ Row key is a time identifier
- ◆ Column names are events
- ◆ Column values are measurements
 - Complex measured values can be serialized (JSON/XML/Thrift/...)
- ◆ Rows can be very wide

Indexing with Tables

- Indexing expresses application intent
- Fast access to specific queries
- Secondary indexes \neq relational indexes
- Use information you have. No pre-reads.

Keyword index

- Use a word as a key
- Columns are the occurrence
- Ex: Index of tag words about videos

```
CREATE TABLE tag_index (  
    tag varchar,  
    videoid uuid,  
    timestamp timestamp,  
    PRIMARY KEY (tag, videoid)  
);
```

Fast



tag	VideoId1	..	VideoIdN
-----	----------	----	----------



Efficient

Partial word index

- Where row size will be large
- Take one part for key, rest for columns name

```
CREATE TABLE email_index (  
    domain varchar,  
    user varchar,  
    username varchar,  
    PRIMARY KEY (domain, user)  
);
```

User: tcodd Email: tcodd@relational.com




```
INSERT INTO email_index (domain, user, username)  
VALUES ('@relational.com', 'tcodd', 'tcodd');
```


Partial word index

- Create partitions + partial indexes

```
CREATE TABLE product_index (  
  store int,  
  part_number0_3 int,  
  part_number4_9 int,  
  count int,  
  PRIMARY KEY ((store,part_number0_3), part_number4_9)  
);
```

Compound row key!



- Store #8675309 has 3 of part# 7079748575

```
INSERT INTO product_index (store,part_number0_3,part_number4_9,count)  
VALUES (8675309,7079,48575,3);
```

```
SELECT count  
FROM product_index  
WHERE store = 8675309  
AND part_number0_3 = 7079  
AND part_number4_9 = 48575;
```

Fast and efficient!



Bit map index – supports ad hoc queries

- Multiple parts to a key
- Create a truth table of the different combinations
- Inserts == the number of combinations
 - 3 fields? 7 options (Not going to use null choice)
 - 4 fields? 15 options
 - $2^n - 1$, where $n = \#$ of dynamic query fields

Bit map index


- Find a car in a lot by variable combinations

Make	Model	Color	Combination
		x	Color
	x		Model
	x	x	Model+Color
x			Make
x		x	Make+Color
x	x		Make+Model
x	x	x	Make+Model+Color

Bit map index - Table create

- Make a table with three different key combos

```
CREATE TABLE car_location_index (  
    make varchar,  
    model varchar,  
    color varchar,  
    vehical_id int,  
    lot_id int,  
    PRIMARY KEY ((make,model,color),vehical_id)  
);
```



Compound row key with three different options

Bit map index - Adding records

- Pre-optimize for 7 possible questions on insert

```
INSERT INTO car_location_index (make,model,color,vehical_id,lot_id)
VALUES ( 'Ford', 'Mustang', 'Blue', 1234, 8675309);
```

```
INSERT INTO car_location_index (make,model,color,vehical_id,lot_id)
VALUES ( 'Ford', 'Mustang', '', 1234, 8675309);
```

```
INSERT INTO car_location_index (make,model,color,vehical_id,lot_id)
VALUES ( 'Ford', '', 'Blue', 1234, 8675309);
```

```
INSERT INTO car_location_index (make,model,color,vehical_id,lot_id)
VALUES ( 'Ford', '', '', 1234, 8675309);
```

```
INSERT INTO car_location_index (make,model,color,vehical_id,lot_id)
VALUES ( '', 'Mustang', 'Blue', 1234, 8675309);
```

```
INSERT INTO car_location_index (make,model,color,vehical_id,lot_id)
VALUES ( '', 'Mustang', '', 1234, 8675309);
```

```
INSERT INTO car_location_index (make,model,color,vehical_id,lot_id)
VALUES ( '', '', 'Blue', 1234, 8675309);
```

Bit map index - Selecting records

- Different combinations now possible

```
SELECT vehical_id,lot_id
FROM car_location_index
WHERE make = 'Ford'
AND model = ''
AND color = 'Blue';
```



vehical_id	lot_id
1234	8675309

```
SELECT vehical_id,lot_id
FROM car_location_index
WHERE make = ''
AND model = ''
AND color = 'Blue';
```



vehical_id	lot_id
1234	8675309
8765	5551212

Requirements

Find orders by customer, supplier, product, and employee

- Optional date range

Need to display order information – header/details

List all employees for a manager

Orders that have not shipped

- By shipper
- By date

Thank you!

Questions?