

# Cassandra Fundamentals

## Module 4

Cassandra access from Java programs



# Module plan

- ◆ C\* Java driver architecture
- ◆ Basic Java API

# Prerequisites

- ◆ JDK (1.7)
- ◆ Eclipse (or other IDE like IntelliJIDEA)
- ◆ Maven
- ◆ set JAVA\_HOME, ECLIPSE\_HOME, M2\_HOME, M2

## C\* Drivers

- ◆ Datastax supported: <http://github.com/datastax>

Driver	Version	2.1 Features	Out of the C* box
Java	2.1.4	Yes	Yes
Python	2.1.4	Yes	Yes
C/C++	1.0.0-rc1	No	No
C#	2.1.2	Yes	No
Ruby	2.0.0	Partial: C* 2.0 API	No
Node.js	1.0.3	No	No


- ◆ ODBC, Clojure, Erlang, Go, Haskell, Perl, PHP, R, Rust, Scala

# Java Driver


- ♦ Reference C\* driver implementation
- ♦ Netty-based (asynchronous)
- ♦ Compression
- ♦ Security (SSL)
- ♦ Automatic fail-over
- ♦ Configurable policies
  - LoadBalancingPolicy, ReconnectionPolicy, RetryPolicy
- ♦ Query tracing support
- ♦ On Maven Central (cassandra-driver-core)
  - depends on Netty, Guava, Metrics

# Request Pipelining

No pipelining



Pipelining



# Configure and Connect

```
Cluster cluster = Cluster.builder()  
    .addContactPoints("127.0.0.1", "192.168.1.100")  
    .setConsistencyLevel(ConsistencyLevel.ONE)  
    .withLoadBalancingPolicy(new DCAwareRoundRobinPolicy("DC1")  
    .withRetryPolicy(DowngradingConsistencyRetryPolicy.INSTANCE)  
    .withReconnectionPolicy(new ConstantReconnectionPolicy(100L))  
    .build();
```

```
Session session = cluster.connect();
```

```
cluster.shutdown();
```

**cassandra.yml**

```
start_native_transport : true  
# native_transport_port : 9042 (default)  
rpc_address : IP address or hostname
```

# Create Keyspace and Table

```
session.execute (“CREATE KEYSPACE myKeyspace WITH replication = ” +  
“{‘class’:’SimpleStrategy’, ‘replication_factor’: 3}”);
```

```
session.execute (“CREATE TABLE myKeyspace.table1 ( ” +  
    “id uuid PRIMARY KEY,” +  
    “col1 integer,” +  
    “col2 text,” +  
    “col3 set<text> );” +  
    );
```



## Write

```
Session session = cluster.connect("myKeyspace");  
  
session.execute(  
    "INSERT INTO table1 (id, col1, col2, col3) " +  
    "VALUES (12345678-1111-2222-3333-1234567890ab, " +  
    "123, 'abc', {'first', 'last'});"  
);  
  
session.execute(  
    "INSERT INTO table1 (col1, col2)  
    VALUES (123, 'abc');"  
);
```

# Read

```
ResultSet results = session.execute(
    "SELECT * FROM table1 " +
    "WHERE id=12345678-1111-2222-3333-1234567890ab or col1=123;"
);

for (Row row: results) {
    String userId = row.getString("id");
    Long col1 = row.getInt("col1");
    String col2 = row.getString("col2");
}
```

# Bound Statement

```
PreparedStatement statement = session.prepare(  
    "SELECT * FROM table WHERE id=? or col1=? or col3=?"  
);  
  
BoundStatement boundStatement = new BoundStatement(statement);  
  
Set<String> stringSet = new HashSet<String>();  
stringSet.add("only");  
stringSet.add("one more");  
  
ResultSet result = session.execute(boundStatement.bind(  
    UUID.fromString("12345678-1111-2222-3333-1234567890ab"),  
    123,  
    stringSet));
```

# Paging Through a ResultSet (Cursor)

```
Statement statement = new SimpleStatement (  
    "SELECT * FROM table WHERE id=1234"  
);  
  
statement.setFetchSize(10);  
  
ResultSet result = session.execute(statement);  
  
Iterator<Row> iter = result.iterator();  
  
while (!result.isFullyFetched()) {  
    result.fetchMoreResults();  
  
    Row row = iter.next();  
  
    System.out.println(row.getString(0));  
  
}
```

# Lightweight Transactions: Success Checking

```
ResultSet rs = session.execute(  
    "INSERT INTO users(username, email, pass) " +  
    "VALUES ('abreiman','abreiman@luxoft.com',12341234ccdd) " +  
    "IF NOT EXISTS");  
  
Row row = rs.one();  
  
row.getBool("applied");
```

# Batch Statement

```
BatchStatement bs = new BatchStatement();  
  
bs.add(new SimpleStatement("INSERT ... VALUES(?)", v1));  
  
bs.add(new SimpleStatement("INSERT ... VALUES(?)", v2));  
  
PreparedStatement ps = session.prepare("INSERT ... VALUES(?)");  
  
bs.add(ps.bind(v1));  
  
bs.add(ps.bind(v2));  
  
Session.execute(bs);
```

# Dynamic Queries with QueryBuilder

```
Statement stmt = QueryBuilder.select()  
    .all()  
    .from("keyspace", "table")  
    .where(eq("id", UUID.fromString("...")));
```

```
Statement stmt = QueryBuilder.insertInto("keyspace","table")  
    .value("id", UUID.fromString("..."))  
    .value("name", "abreiman");
```

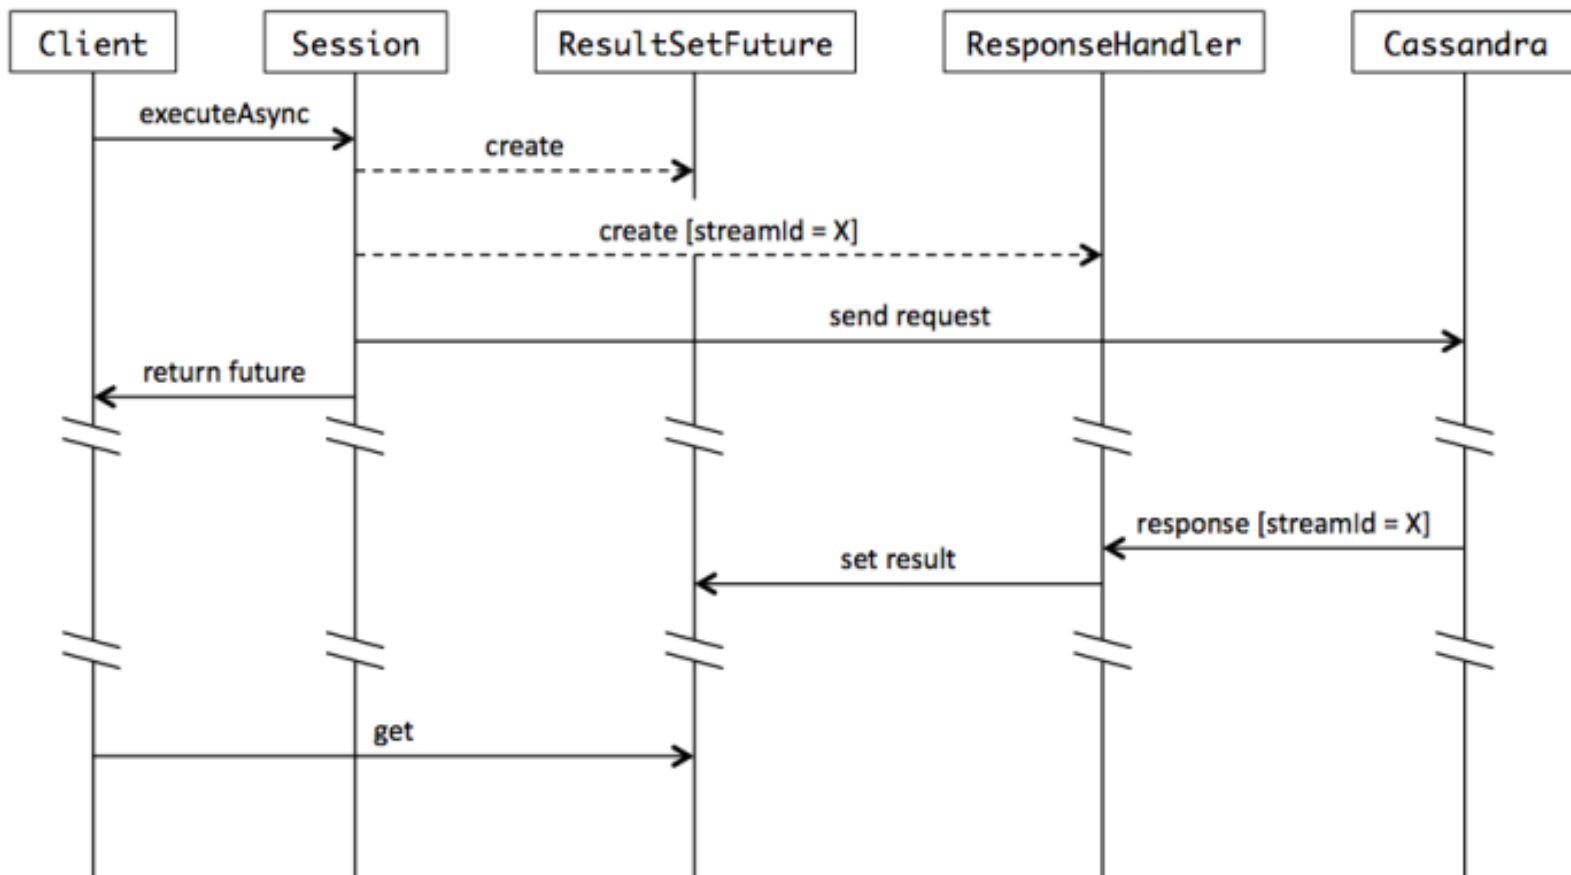
```
Statement stmt = QueryBuilder.update("keyspace","table")  
    .with(set("id", UUID.fromString("..."))  
    .where(eq("id", UUID.fromString("..."))));
```

# Asynchronous I/O

```
PreparedStatement ps = session.prepare("SELECT ... ?");
List<ResultSetFuture> futures = new ArrayList<ResultSetFuture>();
for (int i=0; i<100; i++) {
    ResultSetFuture rsFuture = session.executeAsync(ps.bind(i));
    futures.add(rsFuture);
}
for (ResultSetFuture f: futures) {
    for (Row r: f.getUninterruptibly()) {
        row.getString(0);
    };
}
```



# Asynchronous I/O



# Asynchronous I/O

- ◆ Checking is result ready

```
while (!future.isDone()) {...}
```

- ◆ Blocking wait

```
ResultSet rs = future.get();
```

- ◆ Blocking wait with timeout

```
ResultSet rs = future.get(10, TimeUnit.SECONDS);
```

```
=> catch(TimeoutException e) { future.cancel(true); }
```

## ListenableFuture: Guava Callbacks

```
import com.google.common.util.concurrent.Futures;
import com.google.common.util.concurrent.FutureCallback;
import com.google.common.util.concurrent.MoreExecutors;

ResultSetFuture rsFuture = session.executeAsync("SELECT ...");

Futures.addCallback(rsFuture, new FutureCallback<ResultSet>() {
    @Override public void onSuccess(ResultSet res) { res.get...() }
    @Override public void onFailure(Throwable t) { ... }
},
MoreExecutors.sameThreadExecutor()
);
```

# Cluster Metadata

```
Cluster cluster = ...
```

```
Metadata metadata = cluster.getMetadata();
```

```
String clusterName = metadata.getClusterName();
```

```
for (Host host: metadata.getAllHosts() ) {
```

```
    host.getDatacenter()
```

```
    host.getAddress()
```

```
    host.getRack
```

```
}
```

# Cluster configuration

- ◆ Load Balancing Policies: which node to execute query on?
  - RoundRobinPolicy – round-robin all nodes (as if all nodes are local), first available wins
  - DCAwareRoundRobinPolicy – round-robin local DC nodes, then remote DC nodes
  - TokenAwarePolicy – wrap another policy with token awareness
- ◆ Reconnect Policies: how often to reconnect to a dead node?
  - ConstantReconnectionPolicy
  - ExponentialReconnectionPolicy (default)

# Cluster configuration

- ◆ Retry Policies: what to do in case of node unavailability or request timeout?
  - DefaultRetryPolicy – keeps consistency level while retrying
  - DowngradingConsistencyRetryPolicy – lowers consistency if:
    - ◆ read timeout: 2+ nodes answered but that's not enough
    - ◆ write timeout: 1+ ack with UNLOGGED BATCH
    - ◆ unavailable exception: 1+ alive node
  - FallthroughRetryPolicy – rethrow, never retry
  - LoggingRetryPolicy – logging wrapper

# Exceptions

- ◆ NoHostAvailableException
- ◆ QueryExecutionException
- ◆ QueryValidationException
- ◆ AlreadyExistsException
- ◆ AuthenticationException
- ◆ ...

# Tracing Query Execution

```
Statement statement = new SimpleStatement("SELECT...")  
                        .enableTracing();
```

```
ResultSet result = session.execute(statement);
```

```
ExecutionInfo exInfo = result.getExecutionInfo();
```

```
exInfo.getQueriedHost()
```

```
for (Host host: exInfo.getTriedHosts()) { host.toString() }
```

```
for (QueryTrace.Event ev: exInfo.getQueryTrace().getEvents()) {
```

```
    ev.getDescription(), ev.getSource(), ev.getTimestamp(), ev.getSourceElapsedMicros()
```

```
}
```



# Best Practices

- ◆ Use QUORUM or LOCAL\_QUORUM
- ◆ Use prepared statements
- ◆ Fully qualify tables
- ◆ Use DC and Token Aware policies
- ◆ Use Token Aware policy
- ◆ Use batches of up to 100 inserts
- ◆ CAS contention is at the partition

**Thank you!**

**Questions?**